



Transaction Concept

Comprehensive Course on DBMS For CS & DA

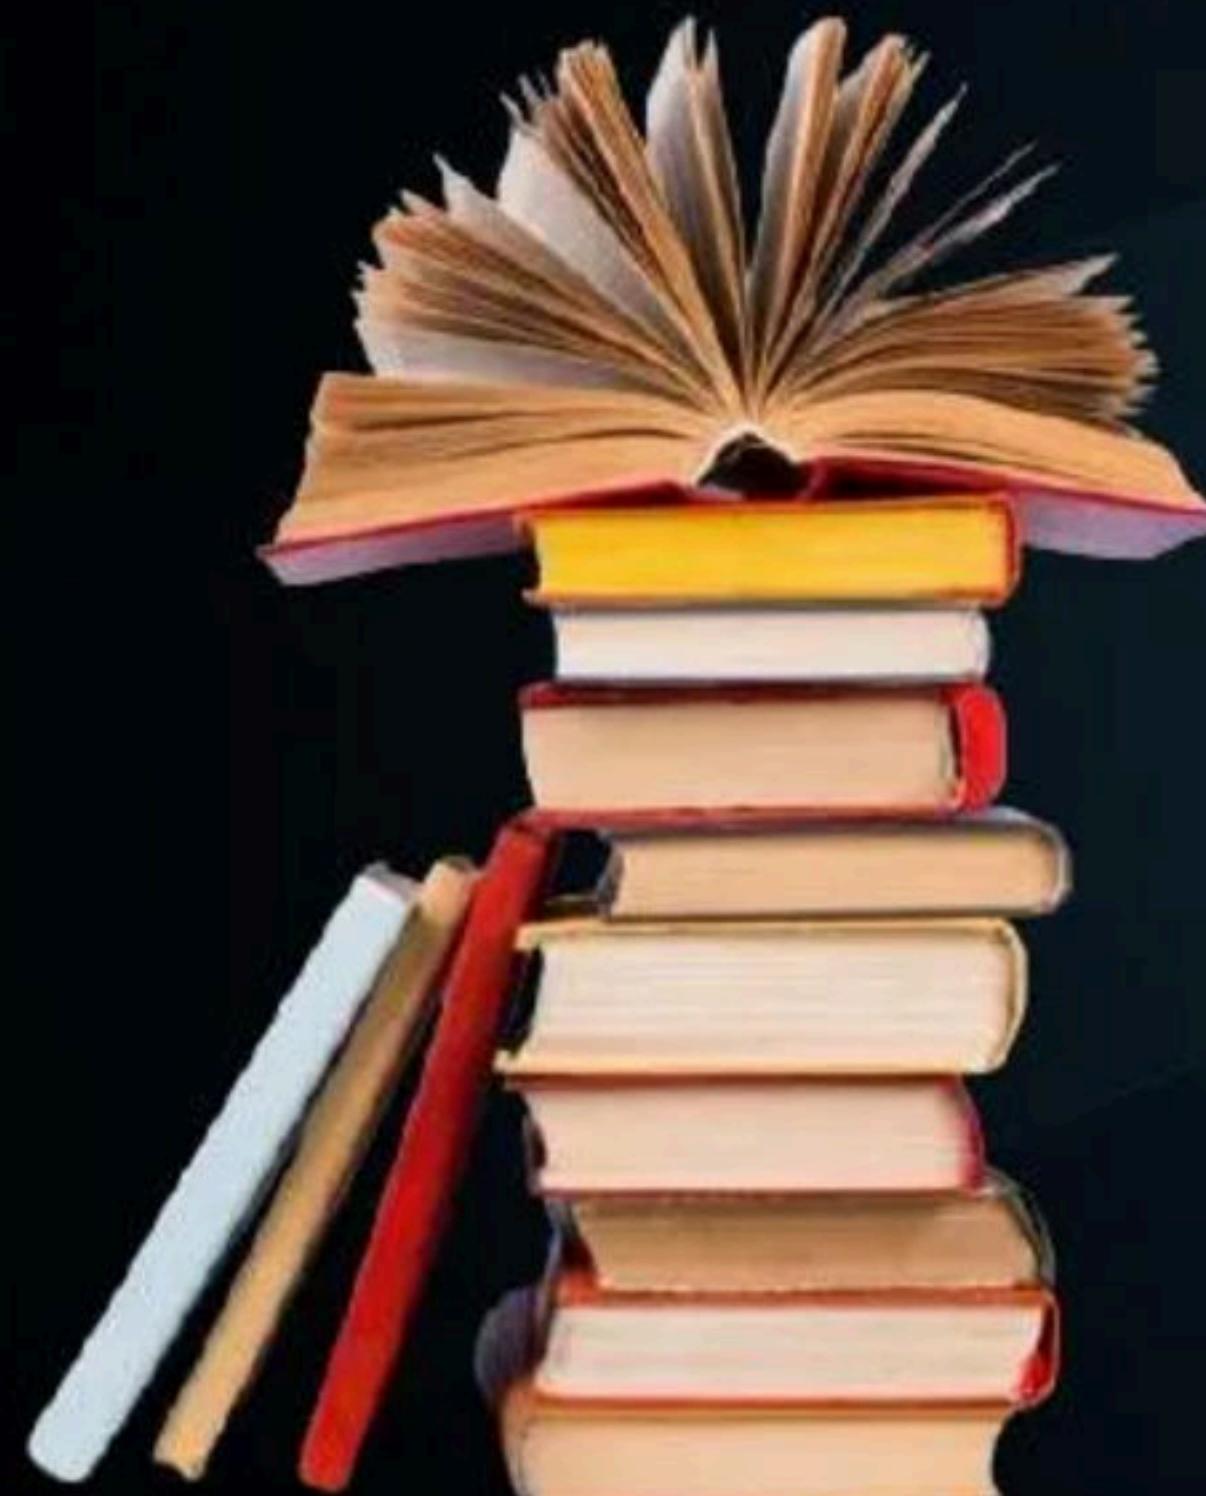
Computer Science & IT Engineering

Transaction Concept



Topics

to be covered

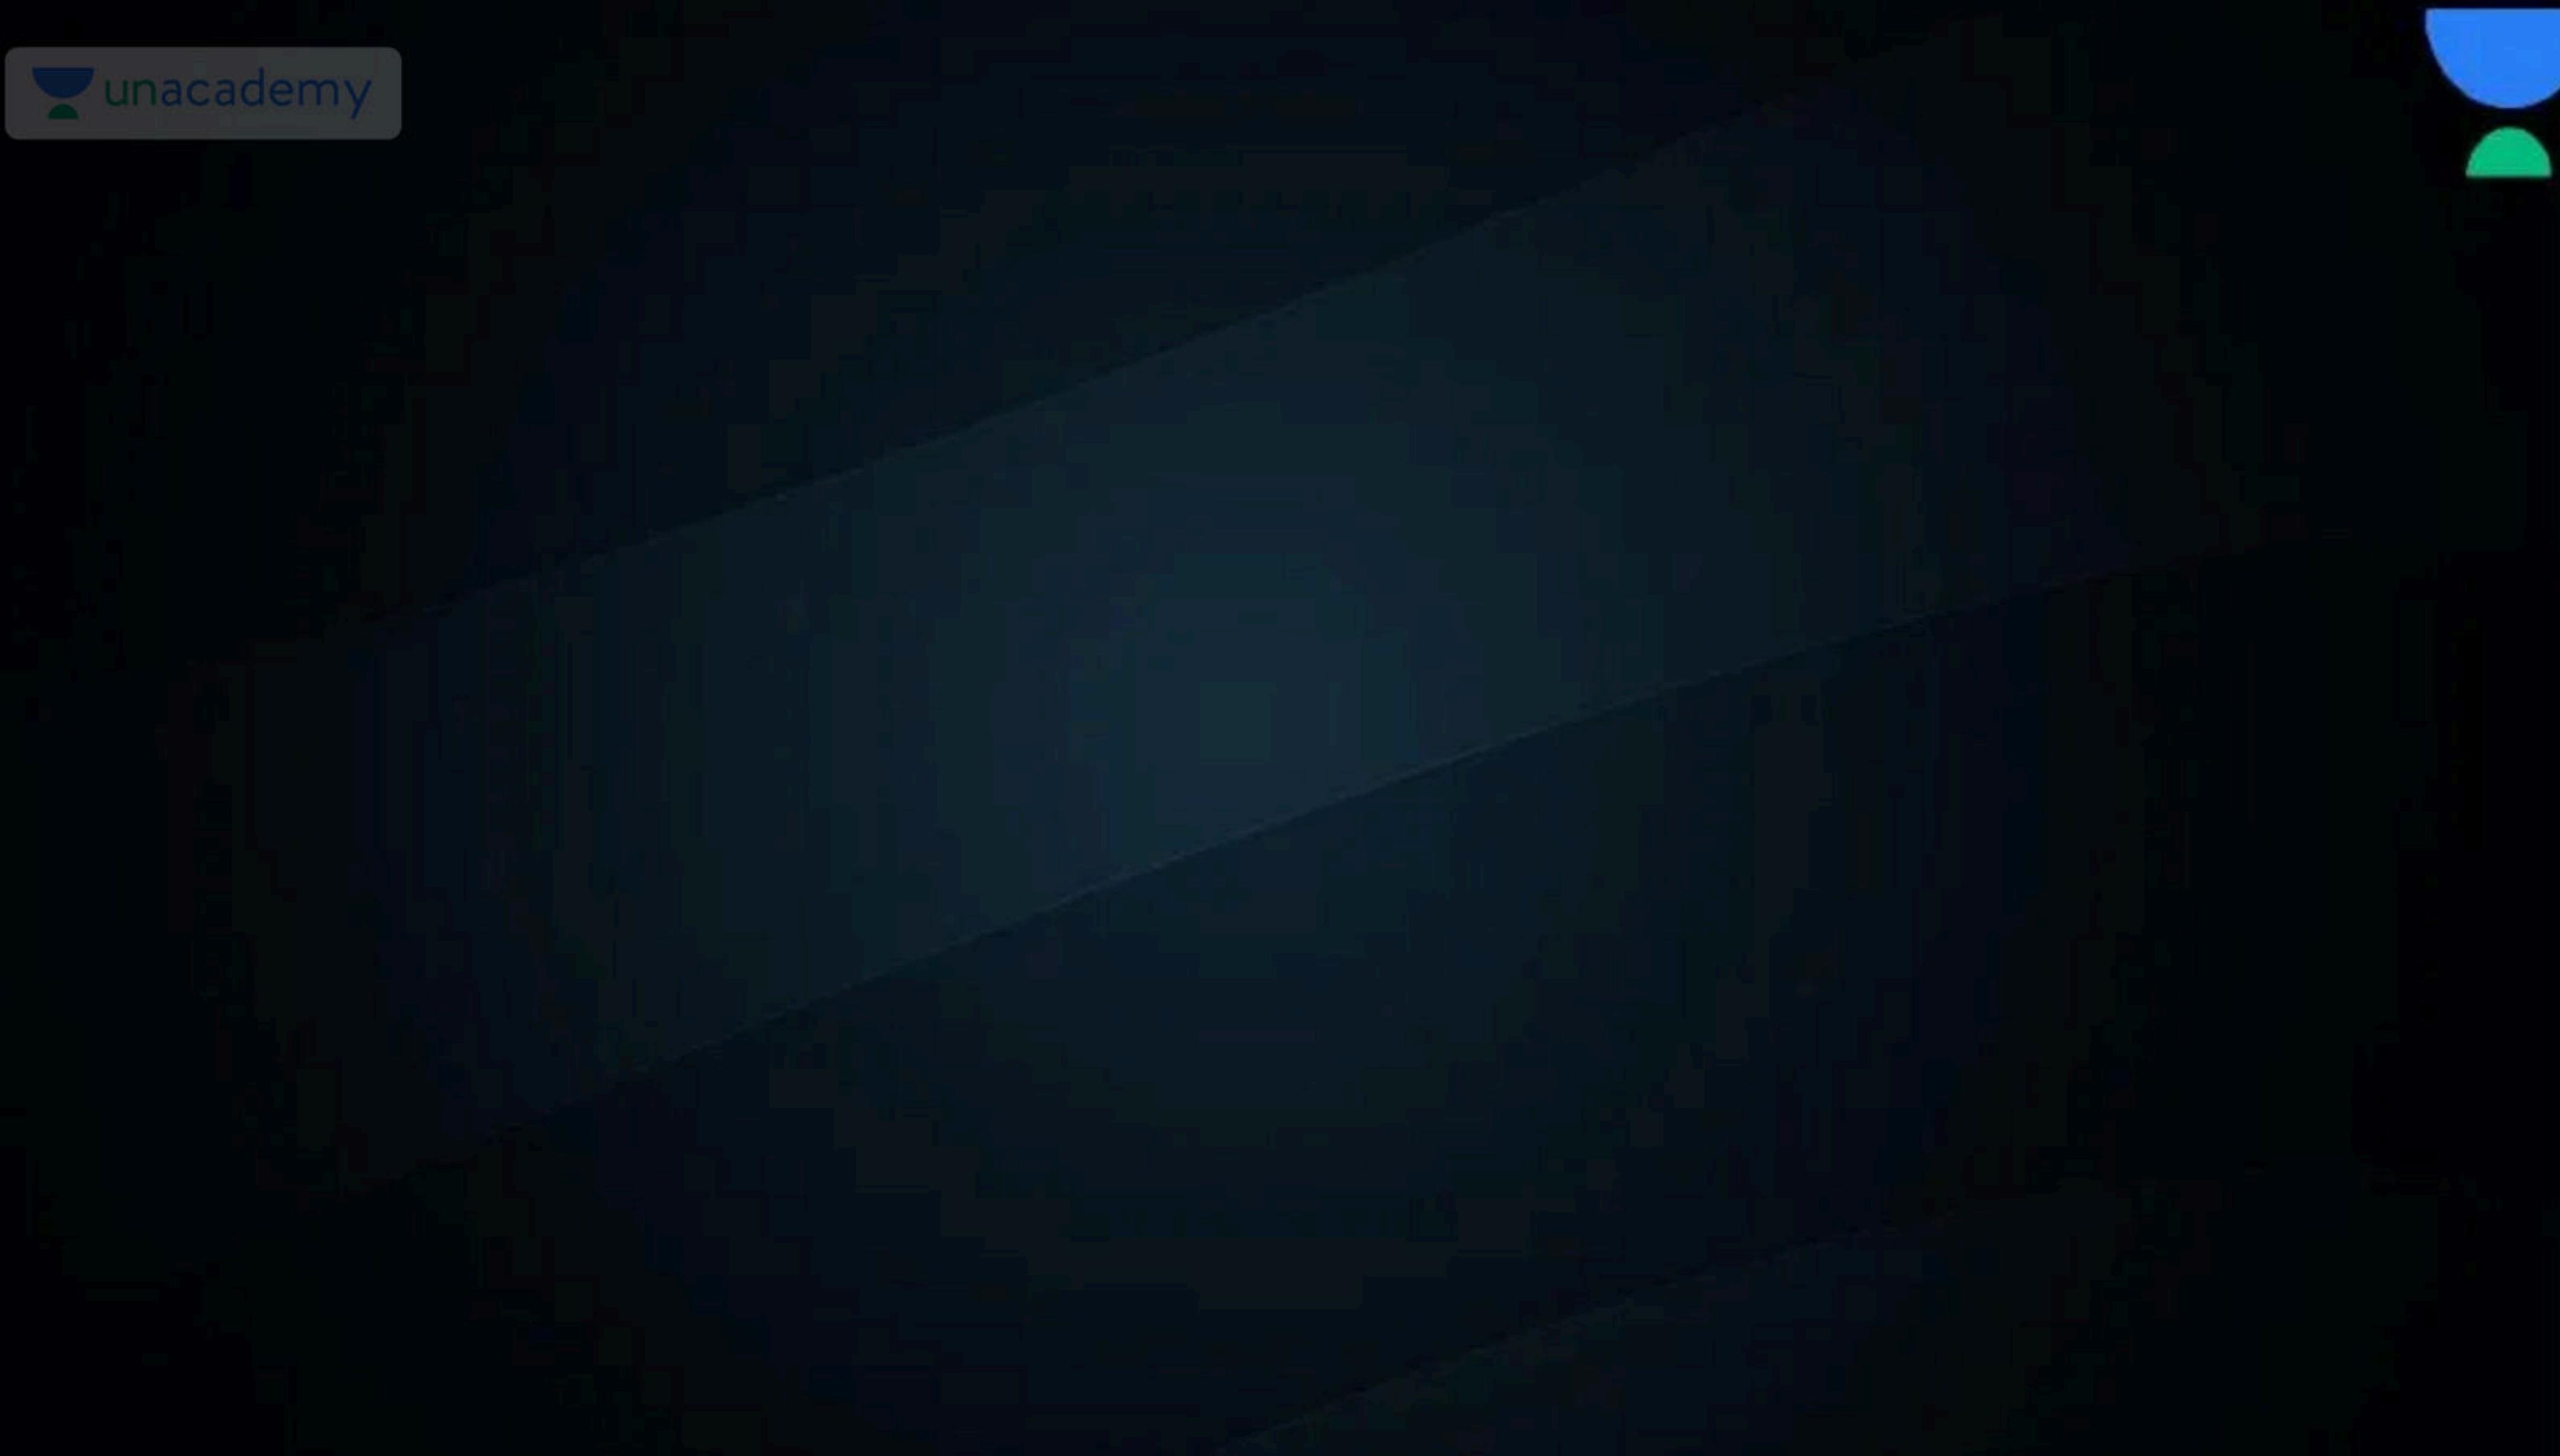
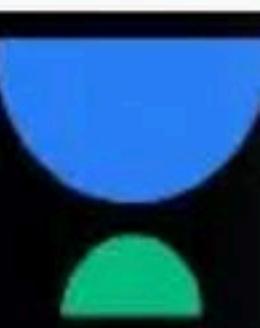


- 1 Transaction Concept
- 2 ACID Properties
- 3 Schedule
- 4 Types of Schedule

Recap of previous Lecture



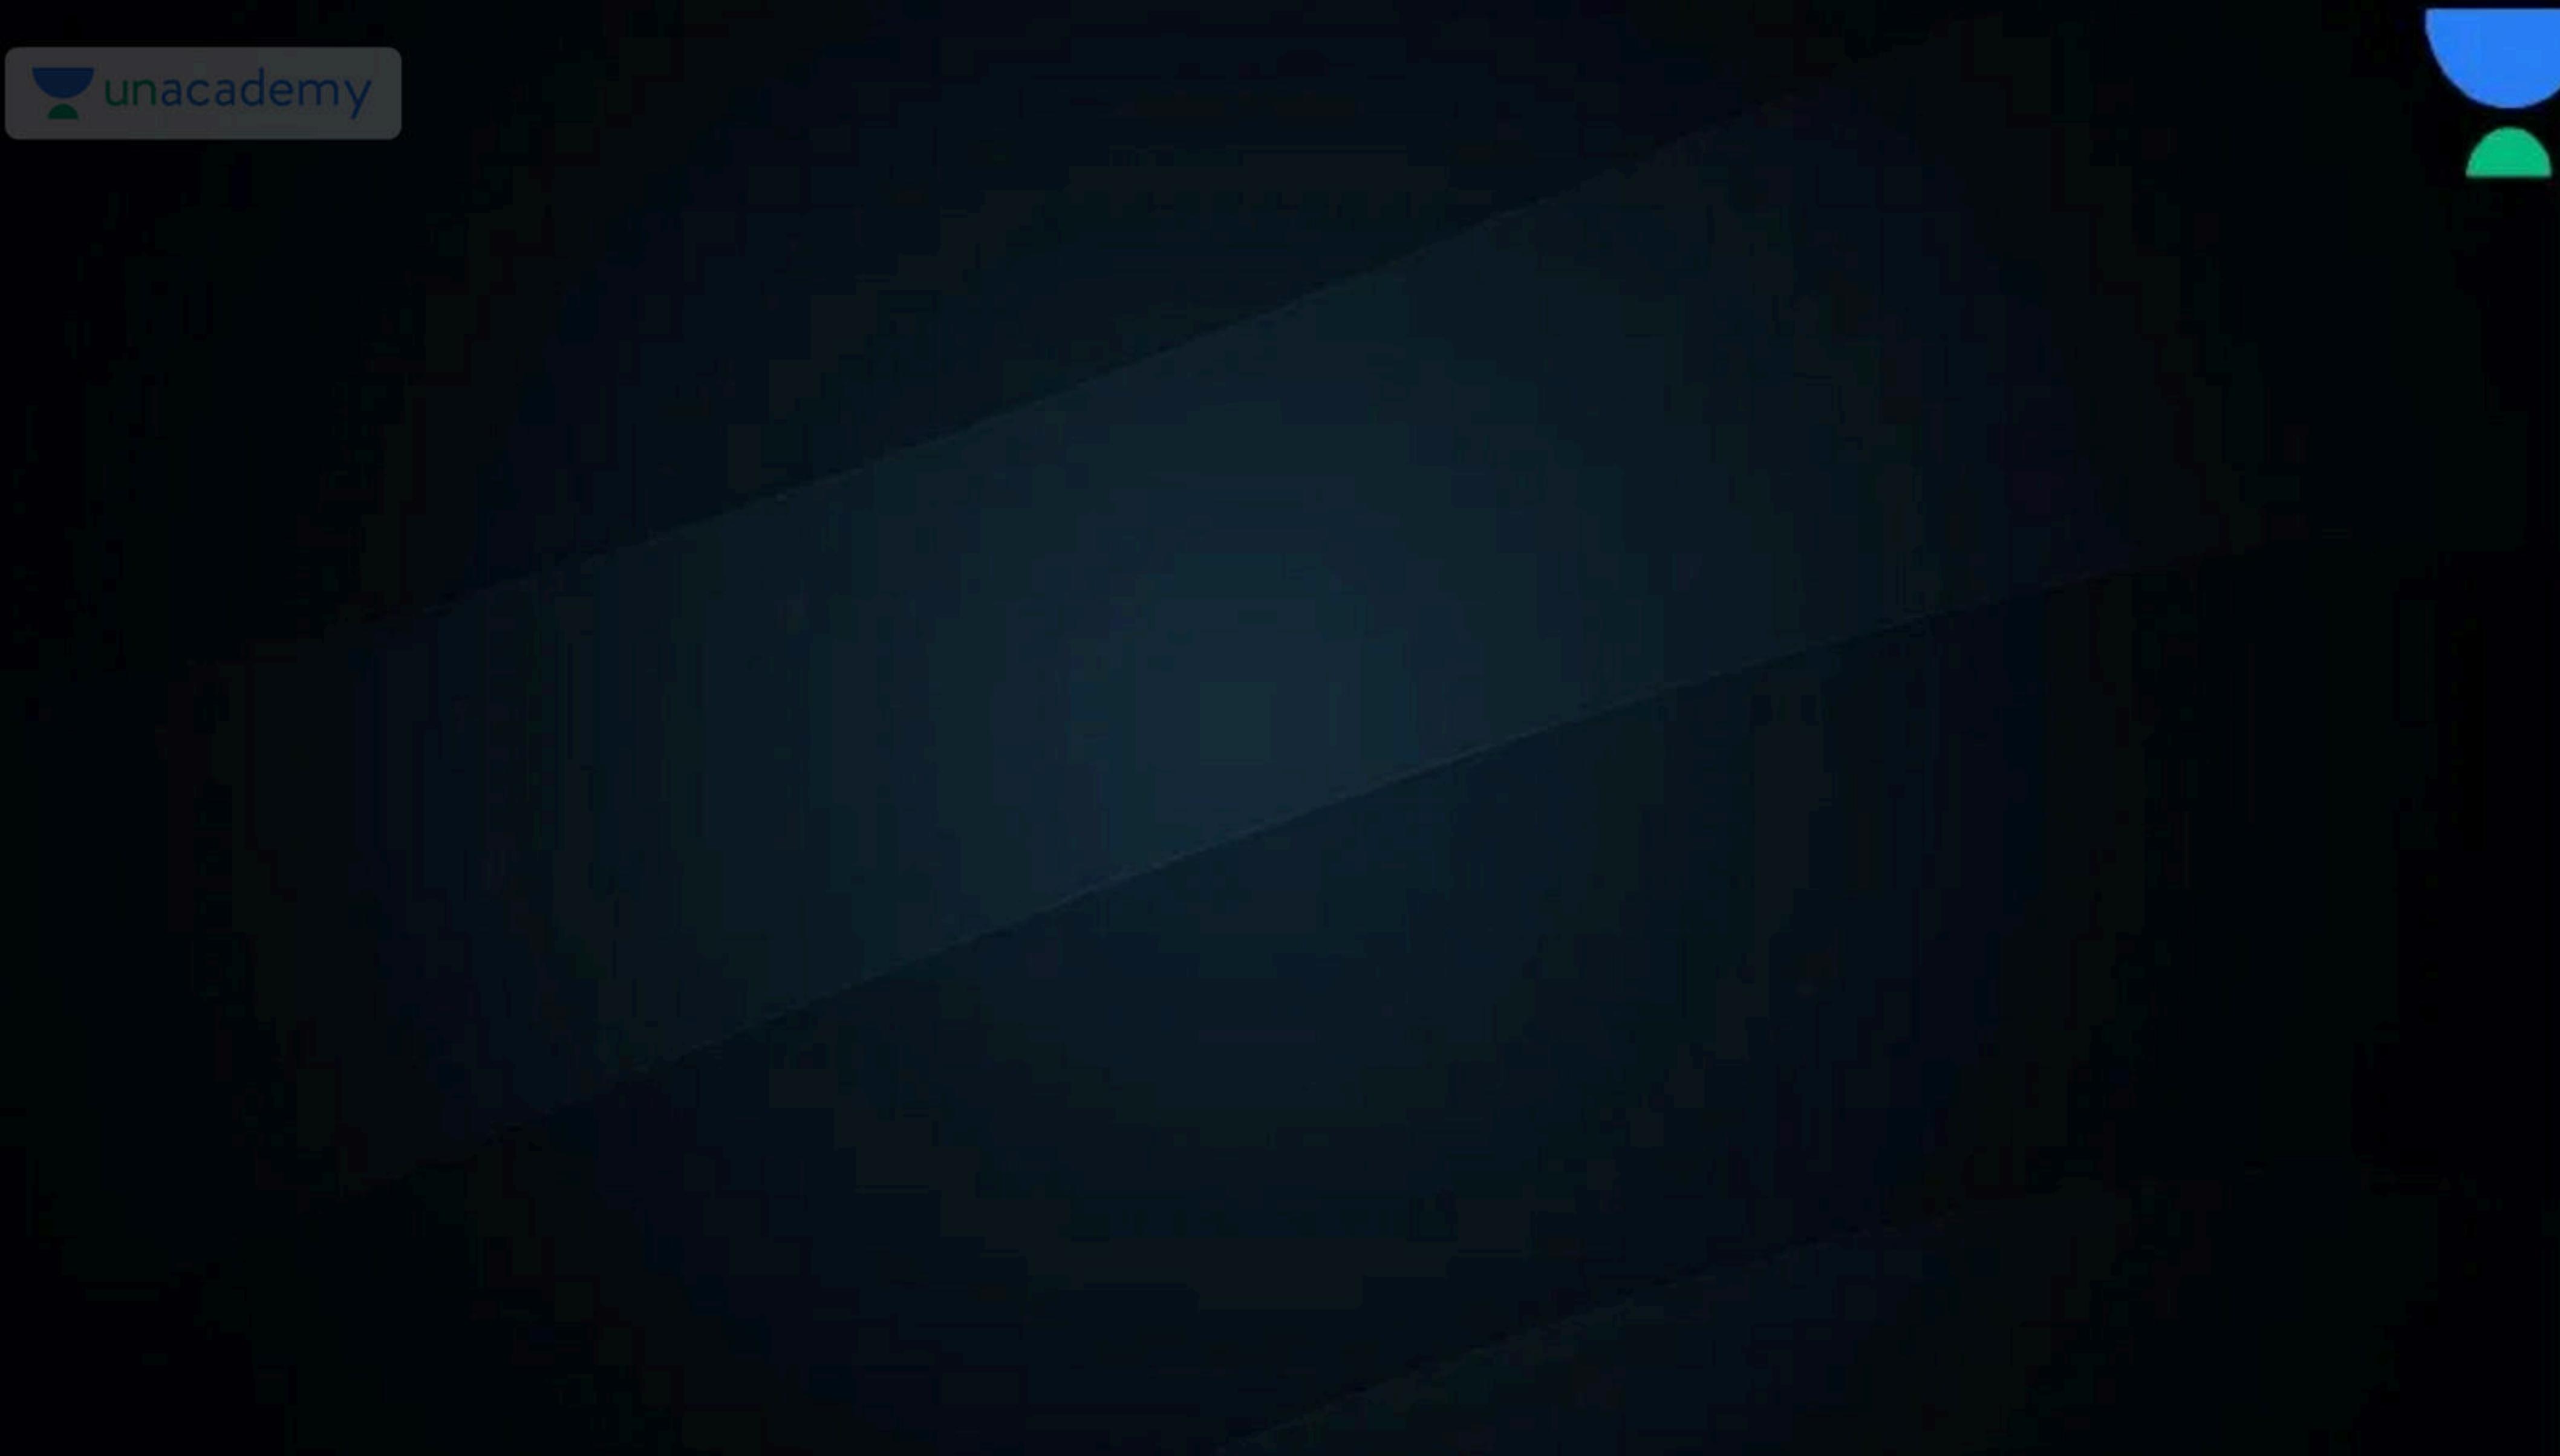
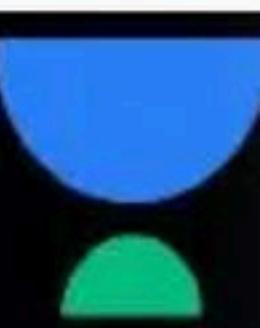
- 1 B+ Tree & GATE PYQ's
- 2 NULL Value Concept & Questions



▲ 1 • Asked by Sunitesh

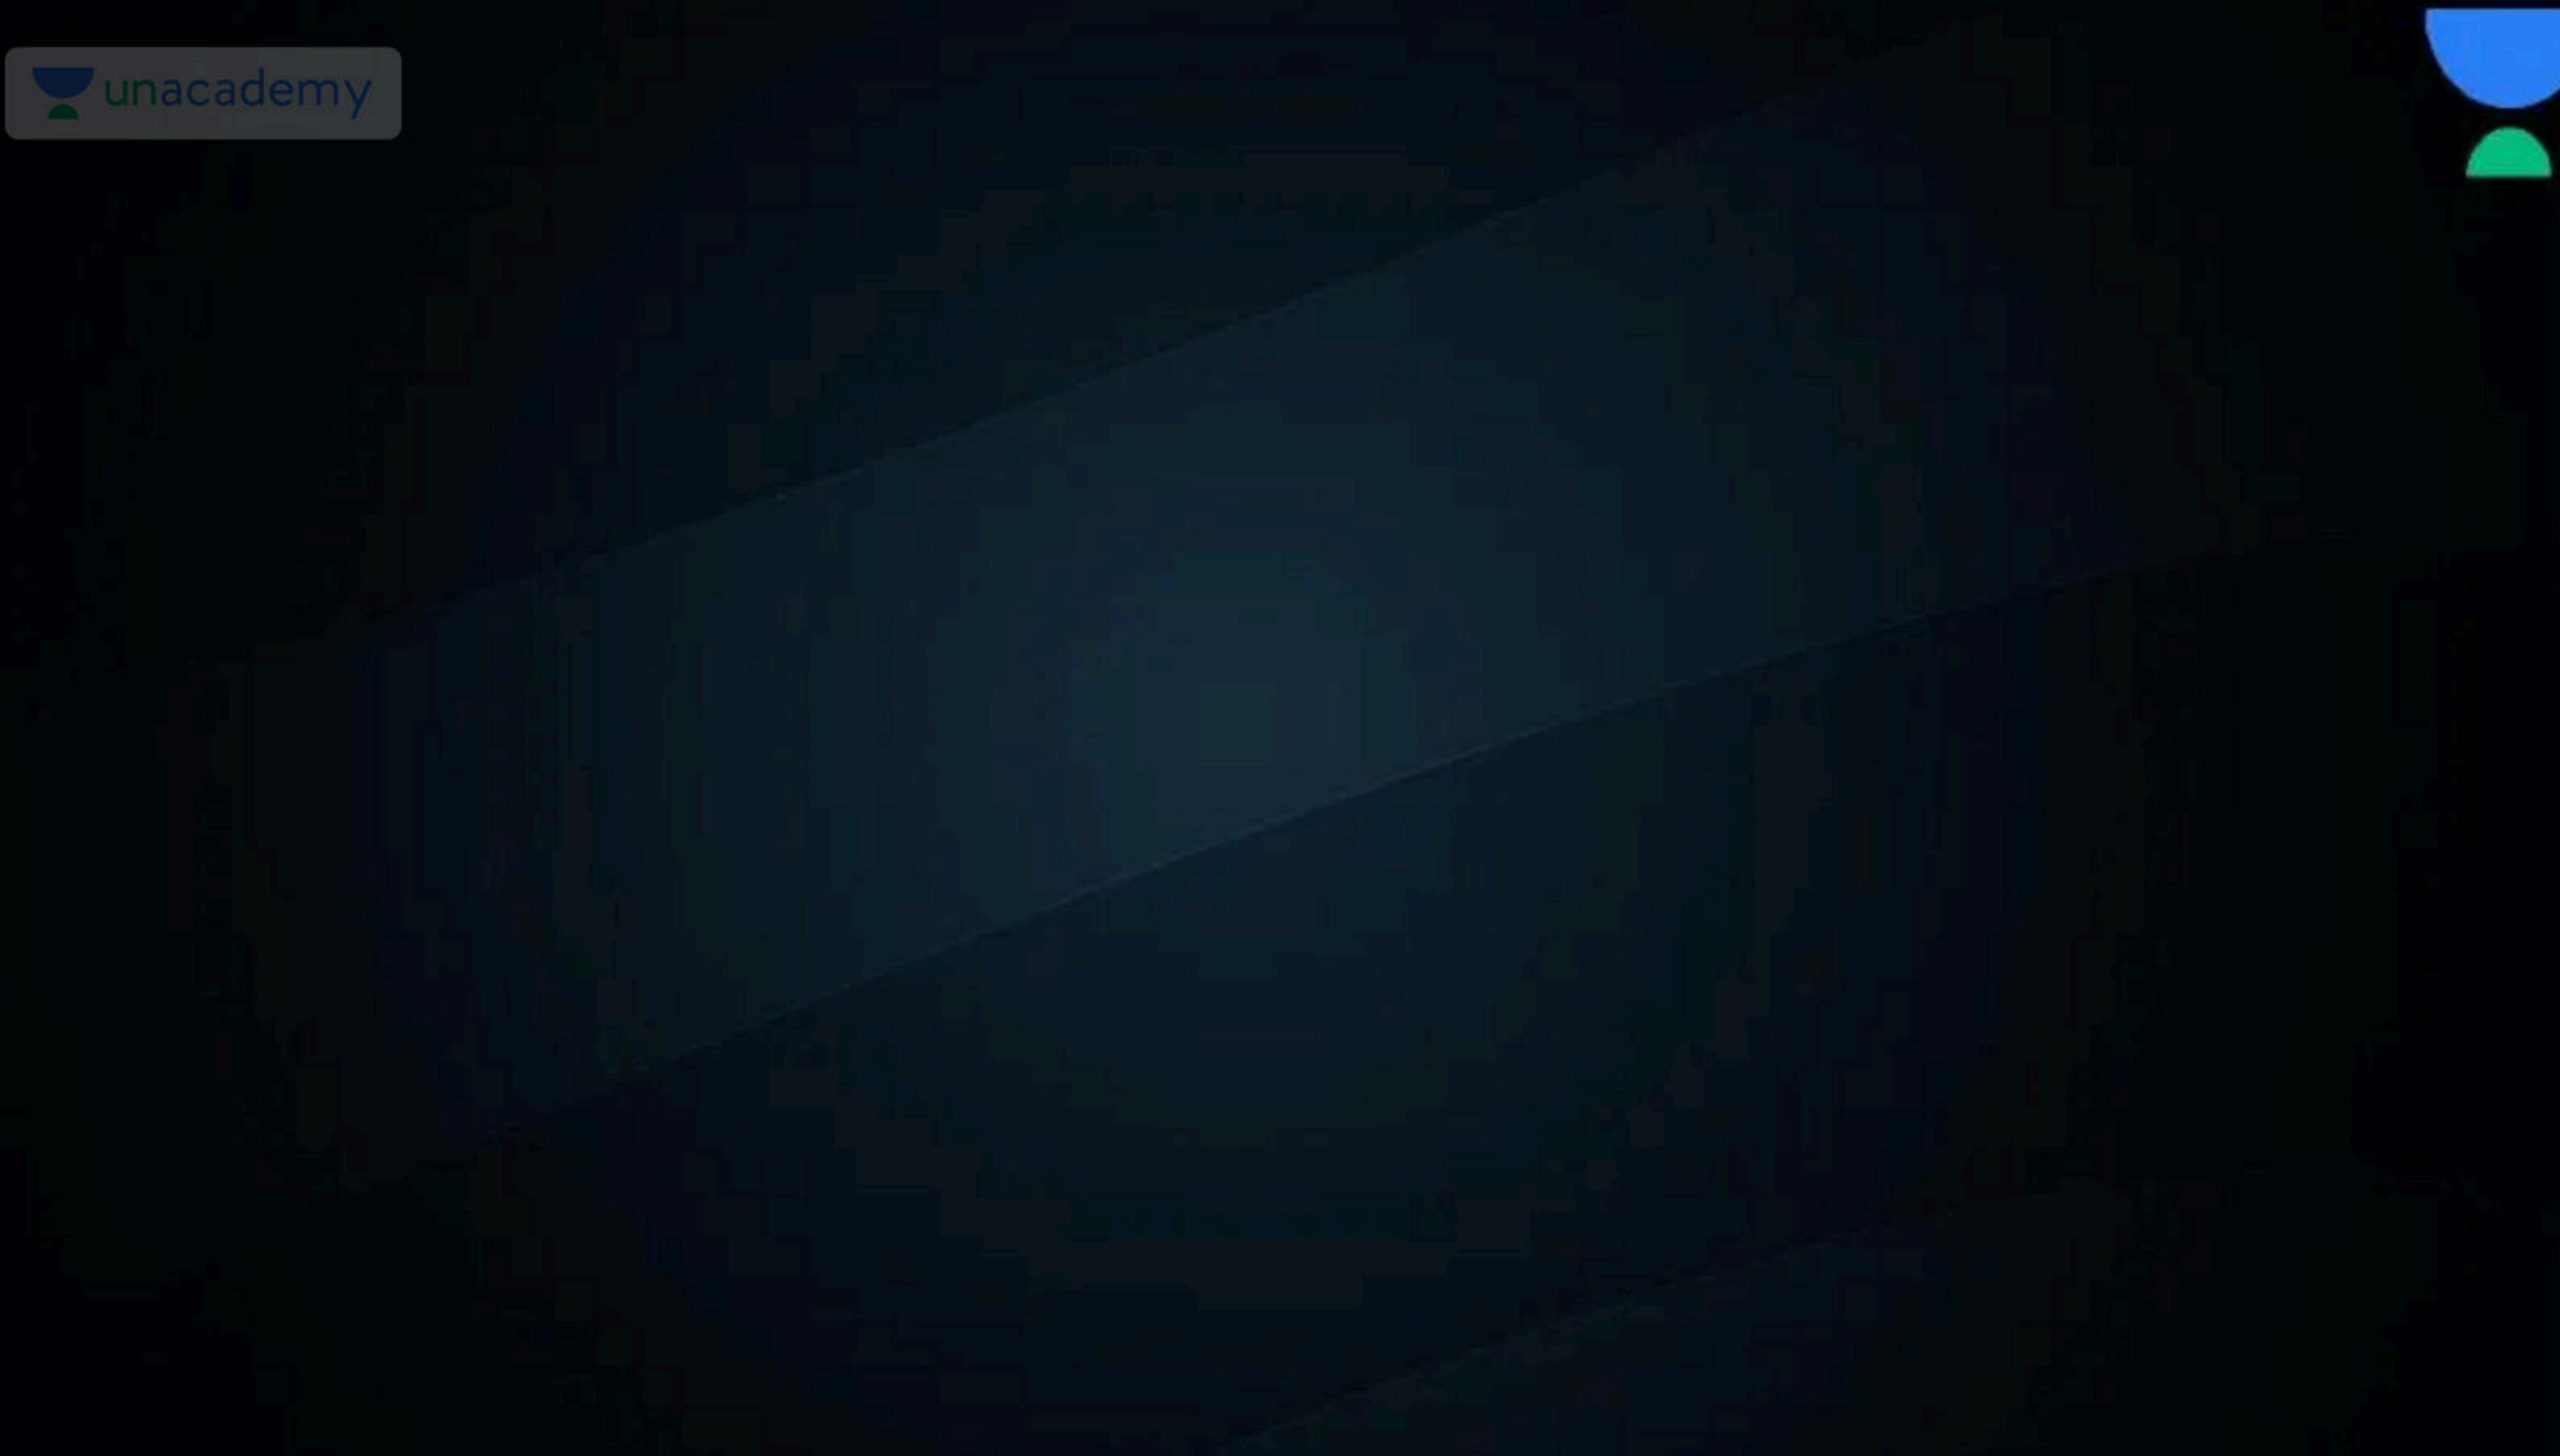
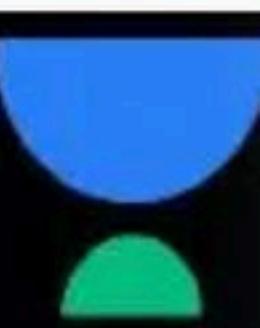
"You are one of those teacher who can make an impact on
any kind of student"

I am greatful to be your student.



Transaction Concept

- A transaction is a unit of program execution that accesses and possibly updates various data items.
- E.g. Transaction to transfer Rs 500 from account A to account B:
 1. read(A)
 2. $A := A - 500$
 3. write(A)
 4. read(B)
 5. $B := B + 500$
 6. write(B)



ACID Properties

- A transaction is a unit of program execution that accesses and possibly updates various data items. To preserve the integrity of data the database system must ensure:
 - A. Atomicity
 - C. Consistency
 - I. Isolation
 - D. Durability

ACID Properties

- **Atomicity:** Either all operations of the transaction are properly reflected in the database or none are.
- **Consistency:** Execution of a transaction in isolation preserves the consistency of the database.
- **Isolation:** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.

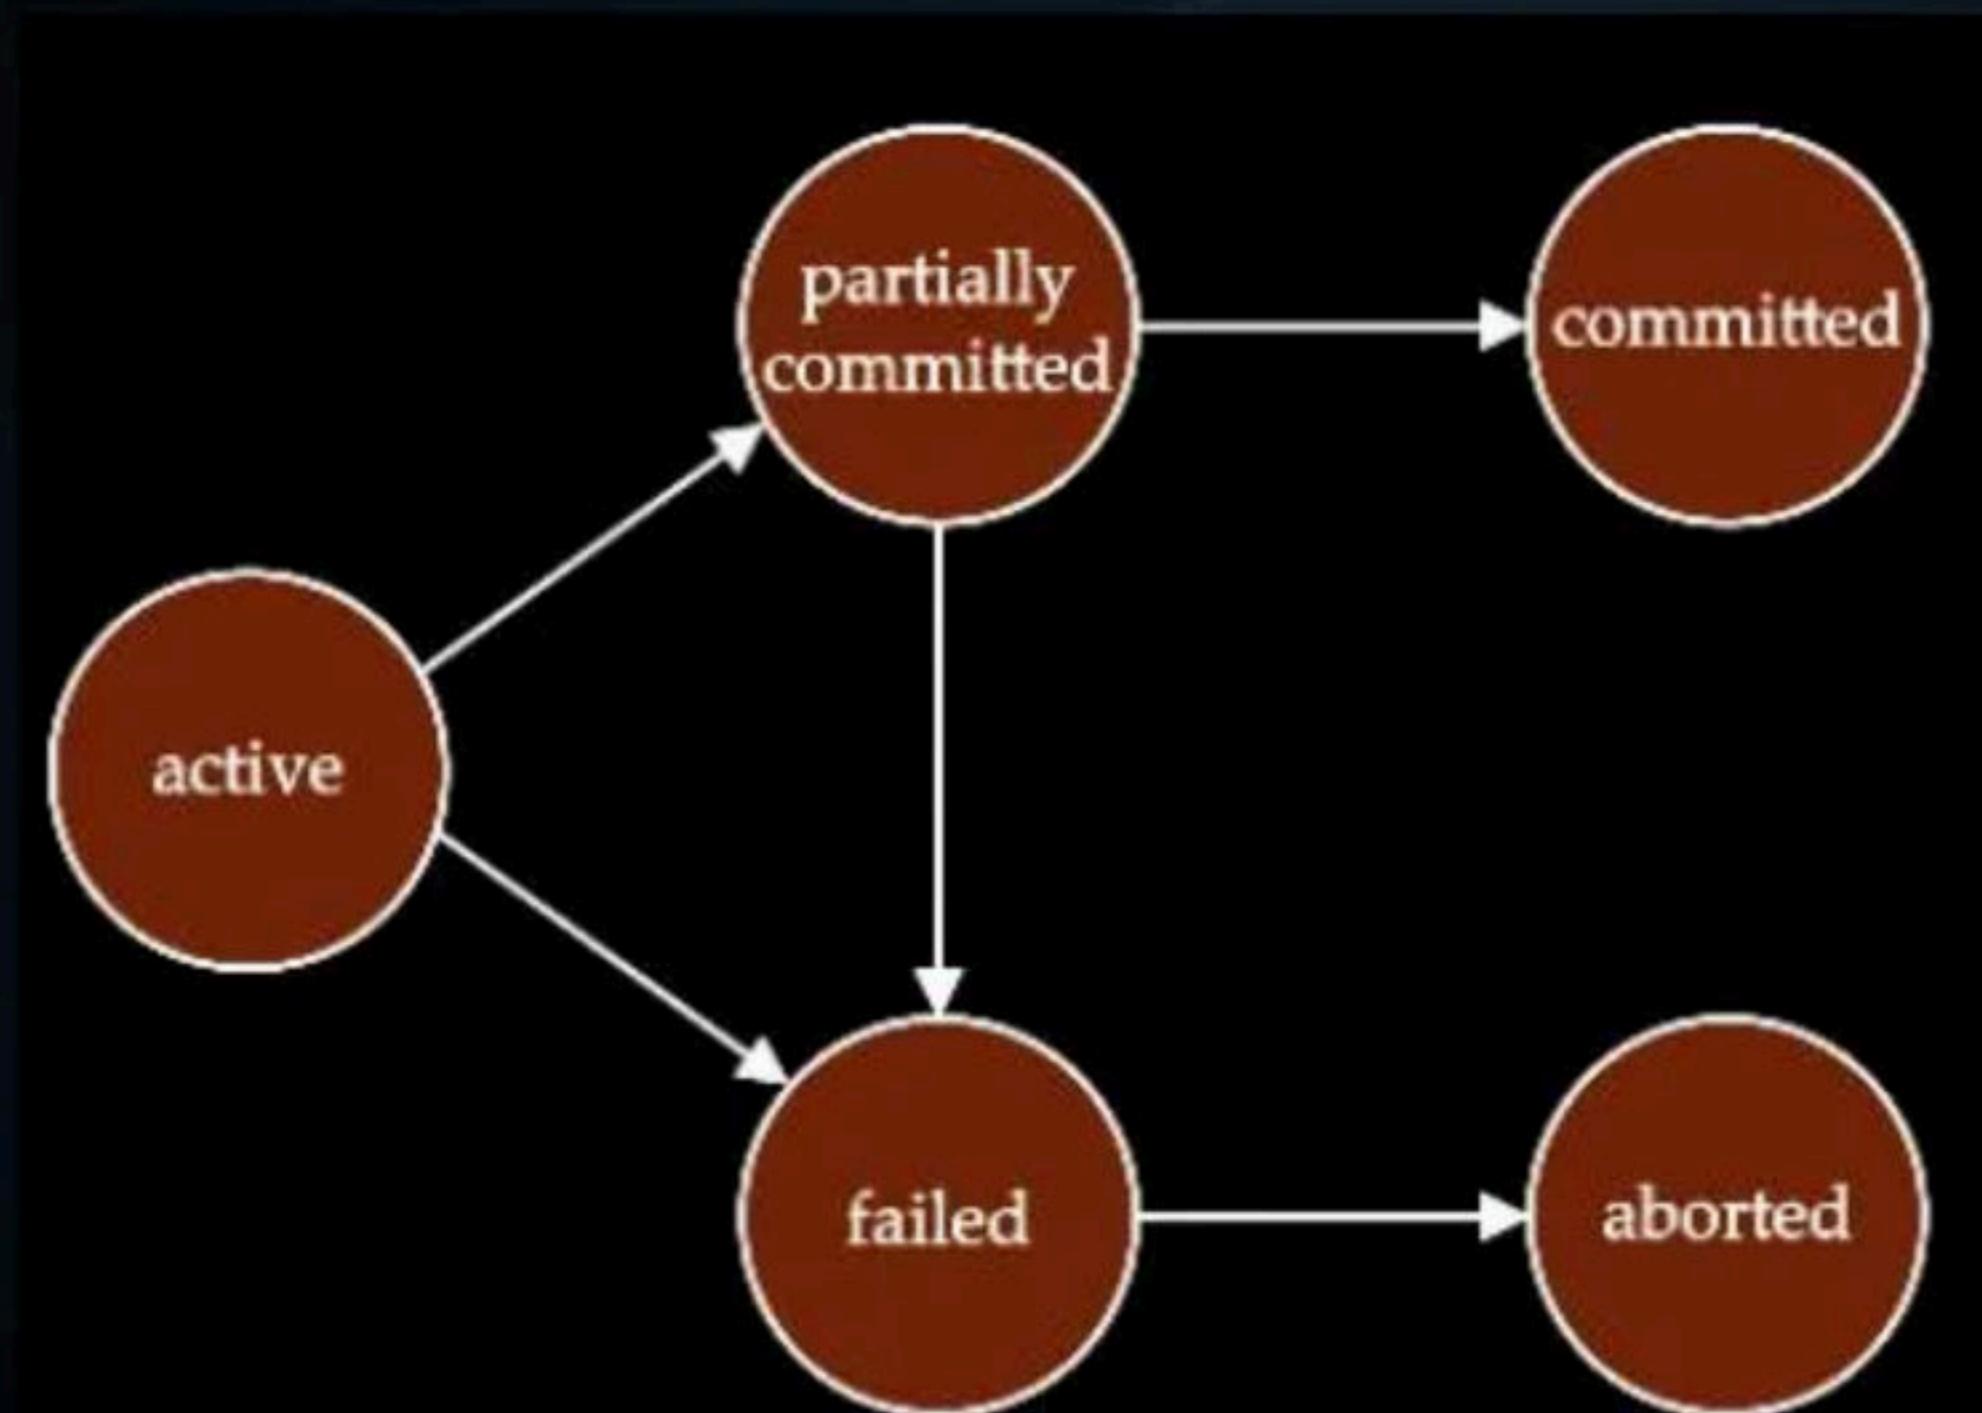
- ❖ That is, for every pair of transactions T_i and T_j , it appears to T_i that either T_j finished execution before T_i started, or T_j started execution after T_i finished.
- **Durability:** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.

Transaction State

- Active : the initial state; the transaction stays in this state while it is executing.
- Partially committed : after the final statement has been executed.
- Failed: after the discovery that normal execution can no longer proceed.
- Aborted: after the transaction has been rolled back and the database restored to its state prior to the start of the transaction.
Two options after it has been aborted:

- ❖ Restart the transaction
 - Can be done only if no internal logical error
- ❖ Kill the transaction
- Committed:** After successful completion.

Transaction State (Cont.)



Schedules

- **Schedule:** a sequences of instructions that specify the chronological order in which instructions of concurrent transactions are executed.
 - ❖ A schedule for a set of transactions must consist of all instructions of those transactions
 - ❖ Must preserve the order in which the instructions appear in each individual transaction.
- A transaction that successfully completes its execution will have a commit instructions as the last statement
 - ❖ By default transaction assumed to execute commit instruction as its last step

- A transaction that fails to successfully complete its execution will have an abort instruction as the last statement.

Let T_1 transfer 100 Rs from A to B, and T_2 transfer 10% of the balance from A to B.

Schedule 1

T_1	T_2
read (A)	
$A := A - 100$	
write (A)	
read (B)	
$B := B + 100$	
write (B)	
commit	

T_1	T_2
	read (A)
	$temp := A * 0.1$
	$A := A - temp$
	write (A)
	read (B)
	$B := B + temp$
	write (B)
	Commit

$S_1 < T_1 \text{ } T_2 >$

Serial schedule in which T_1 is followed by T_2 :

Schedule 2

T_1	T_2
read (A)	
$A := A - 100$	
write (A)	
read (B)	
$B := B + 100$	
write (B)	
commit	

T_1	T_2
	read (A)
	$A := A - 100$
	write (A)
	read (B)
	$B := B + 100$
	write (B)
	commit

$S_2 < T_2 \text{ } T_1 >$

serial schedule where T_2 is followed by T_1

Schedule 3**T₁**

read (A)
 $A := A - 100$
 write (A)

T₂

read (A)
 $temp := A * 0.1$
 $A := A - temp$
 write (A)

read (B)
 $B := B + 100$
 write (B)
 commit

read (B)
 $B := B + temp$
 write (B)
 Commit

C₁**Schedule 4****T₁**

read (A)
 $A := A - 100$

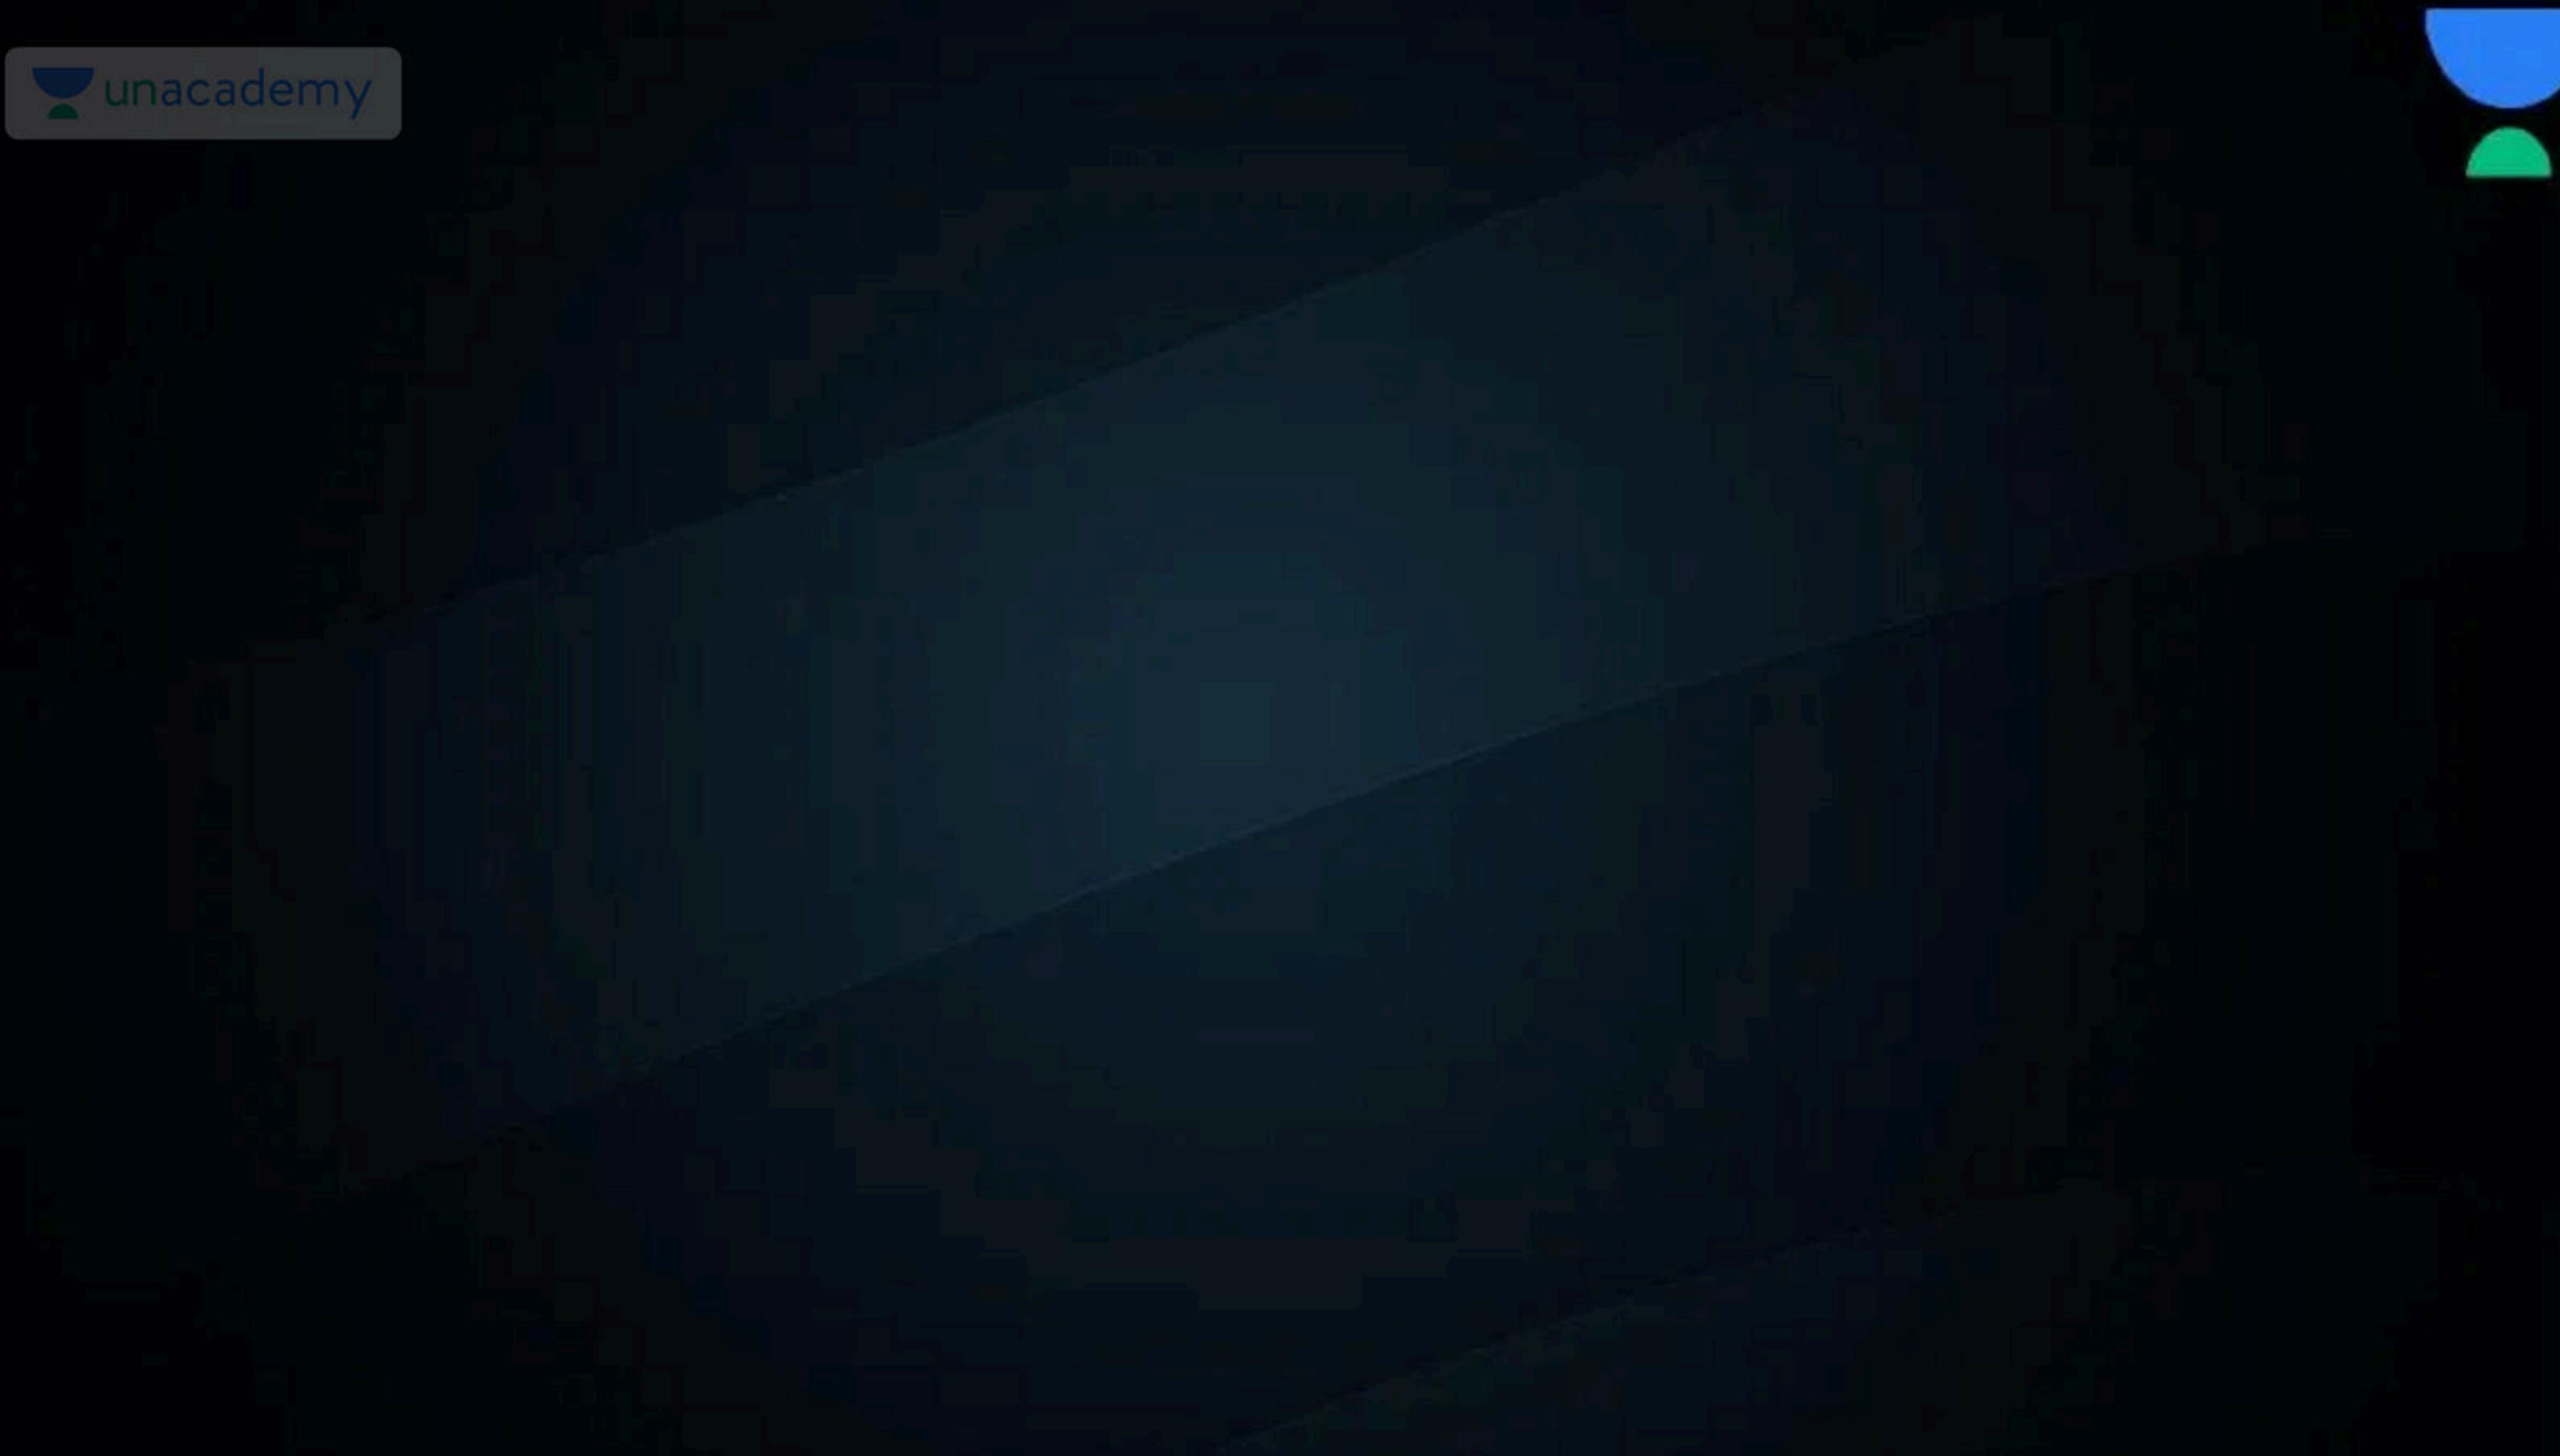
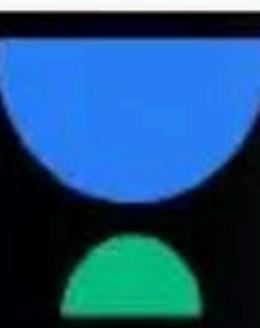
T₂

read (A)
 $temp := A * 0.1$
 $A := A - temp$
 write (A)

write (A)
 read (B)
 $B := B + 100$
 write (B)
 commit

read (B)
 $B := B + temp$
 write (B)
 Commit

C₂



Serial Schedule

- ❑ After Commit of one transaction, begins (Start) another transaction.
- ❑ Number of possible serial Schedules with ‘n’ transactions is “ $n!$ ”
- ❑ The execution sequence of Serial Schedule always generates consistent result.

Example

S : R₁(A) W₁(A) Commit (T₁) R₂(A)W₂(A) commit (T₂).

- Serial Schedule always produce correct result (integrity guaranteed) as no resource sharing.

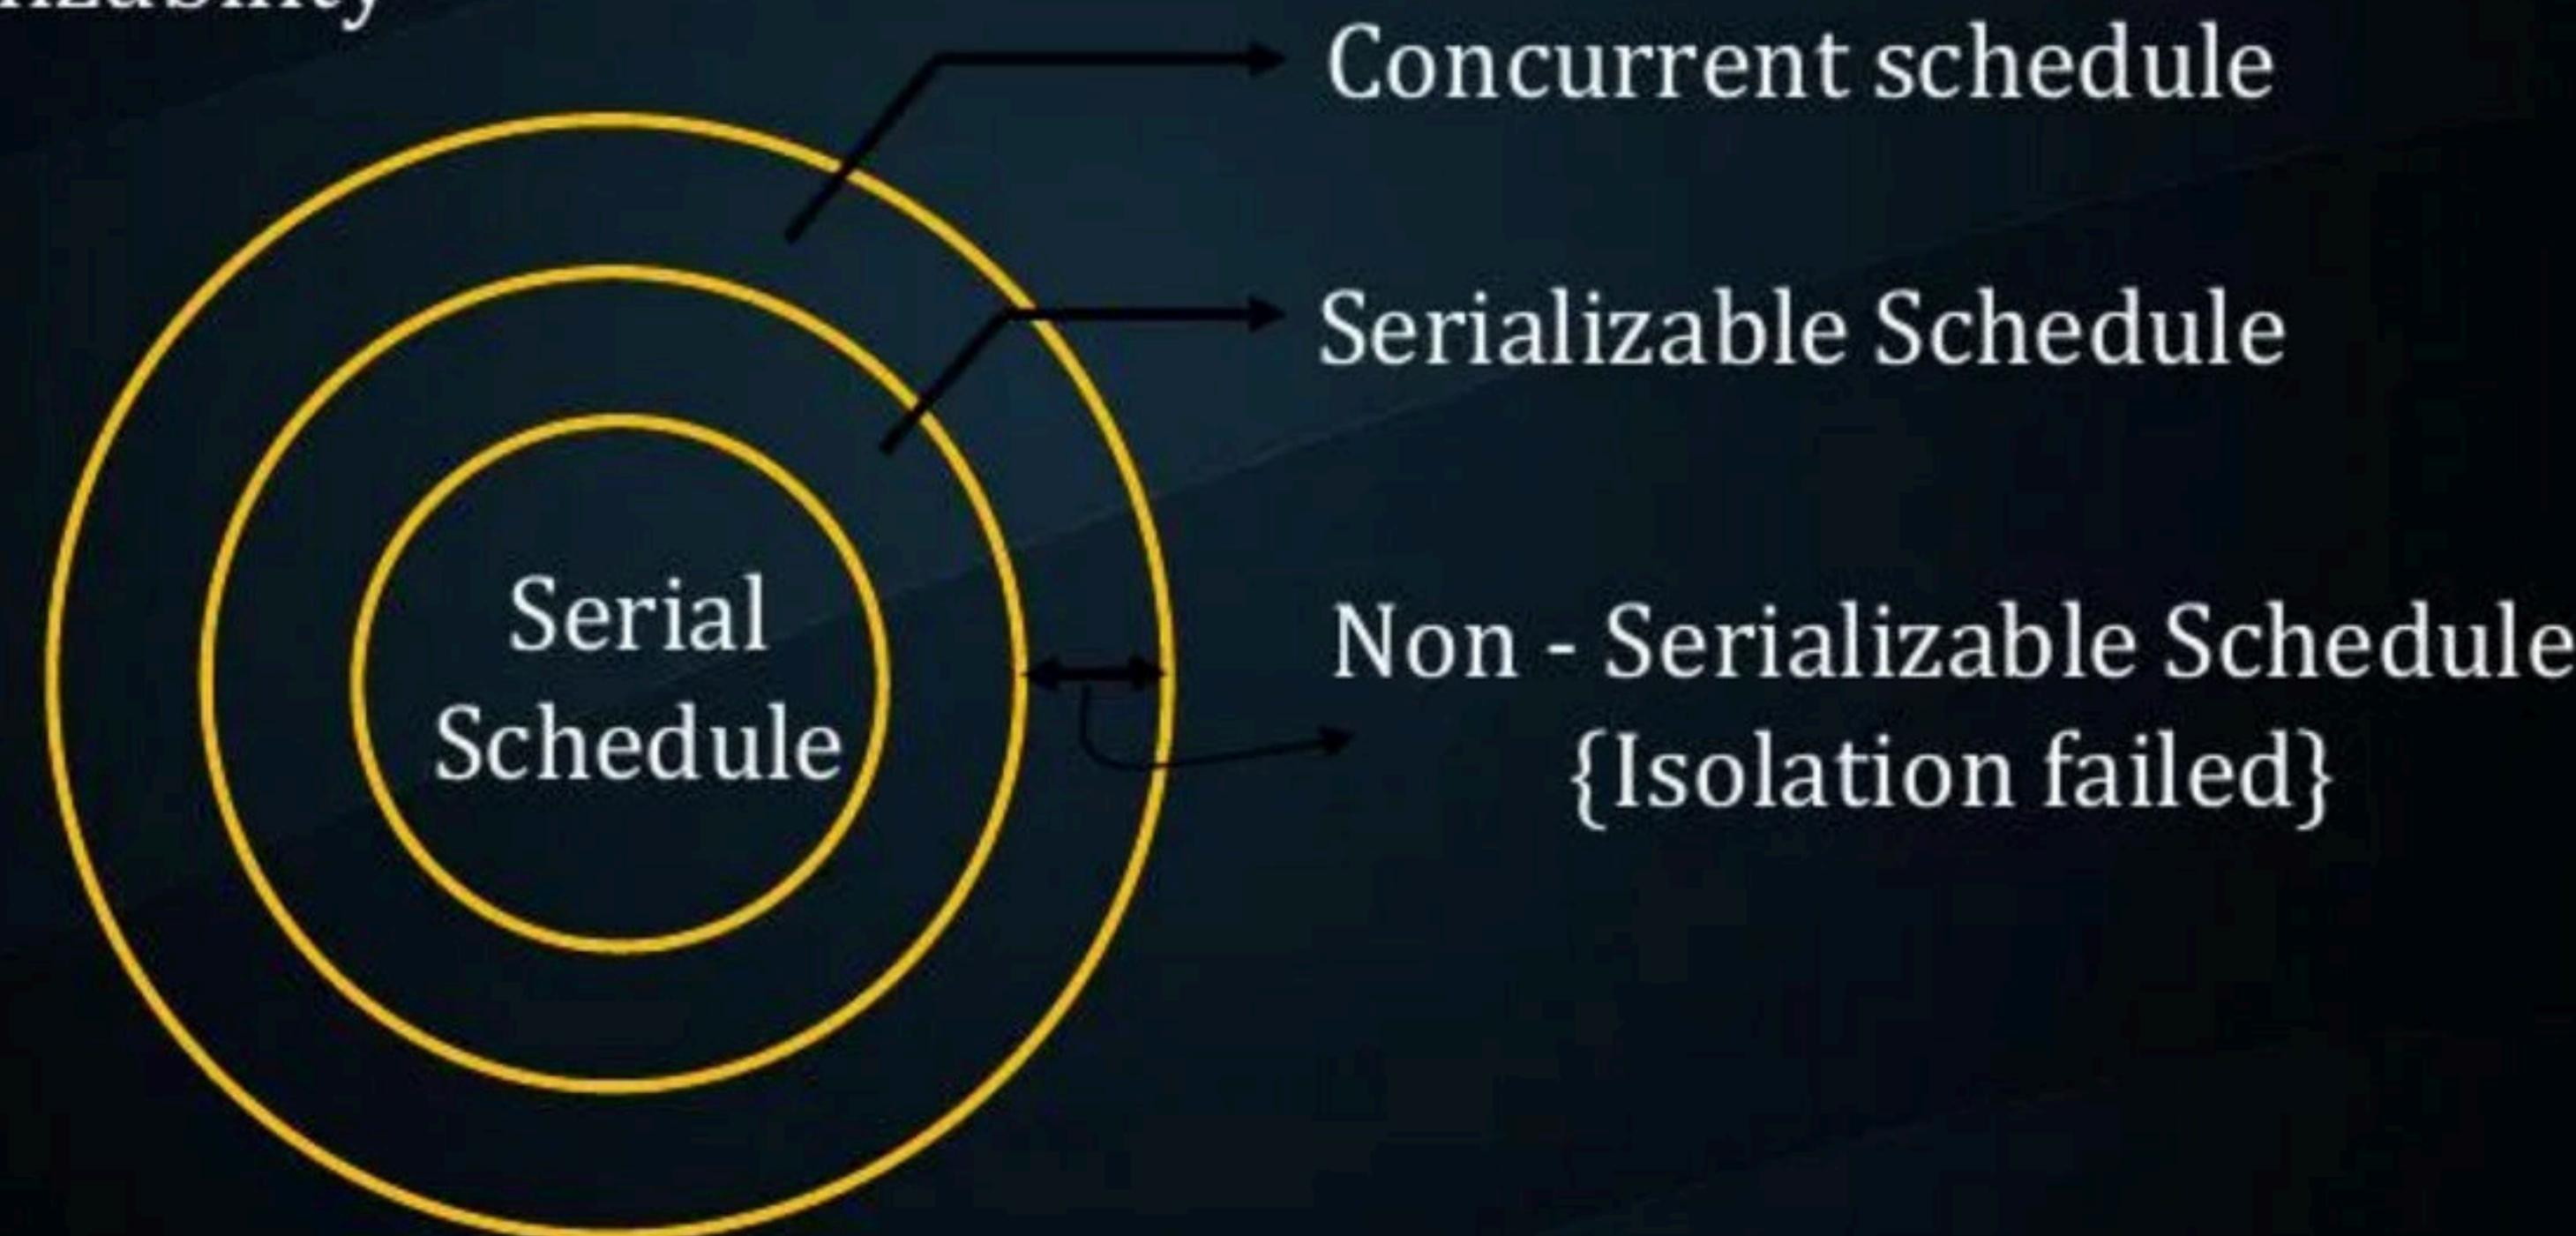
Disadvantage

- Less degree of concurrency.
- Through put of system is low.
- It allows transactions to execute one after another.

Serializable Schedule

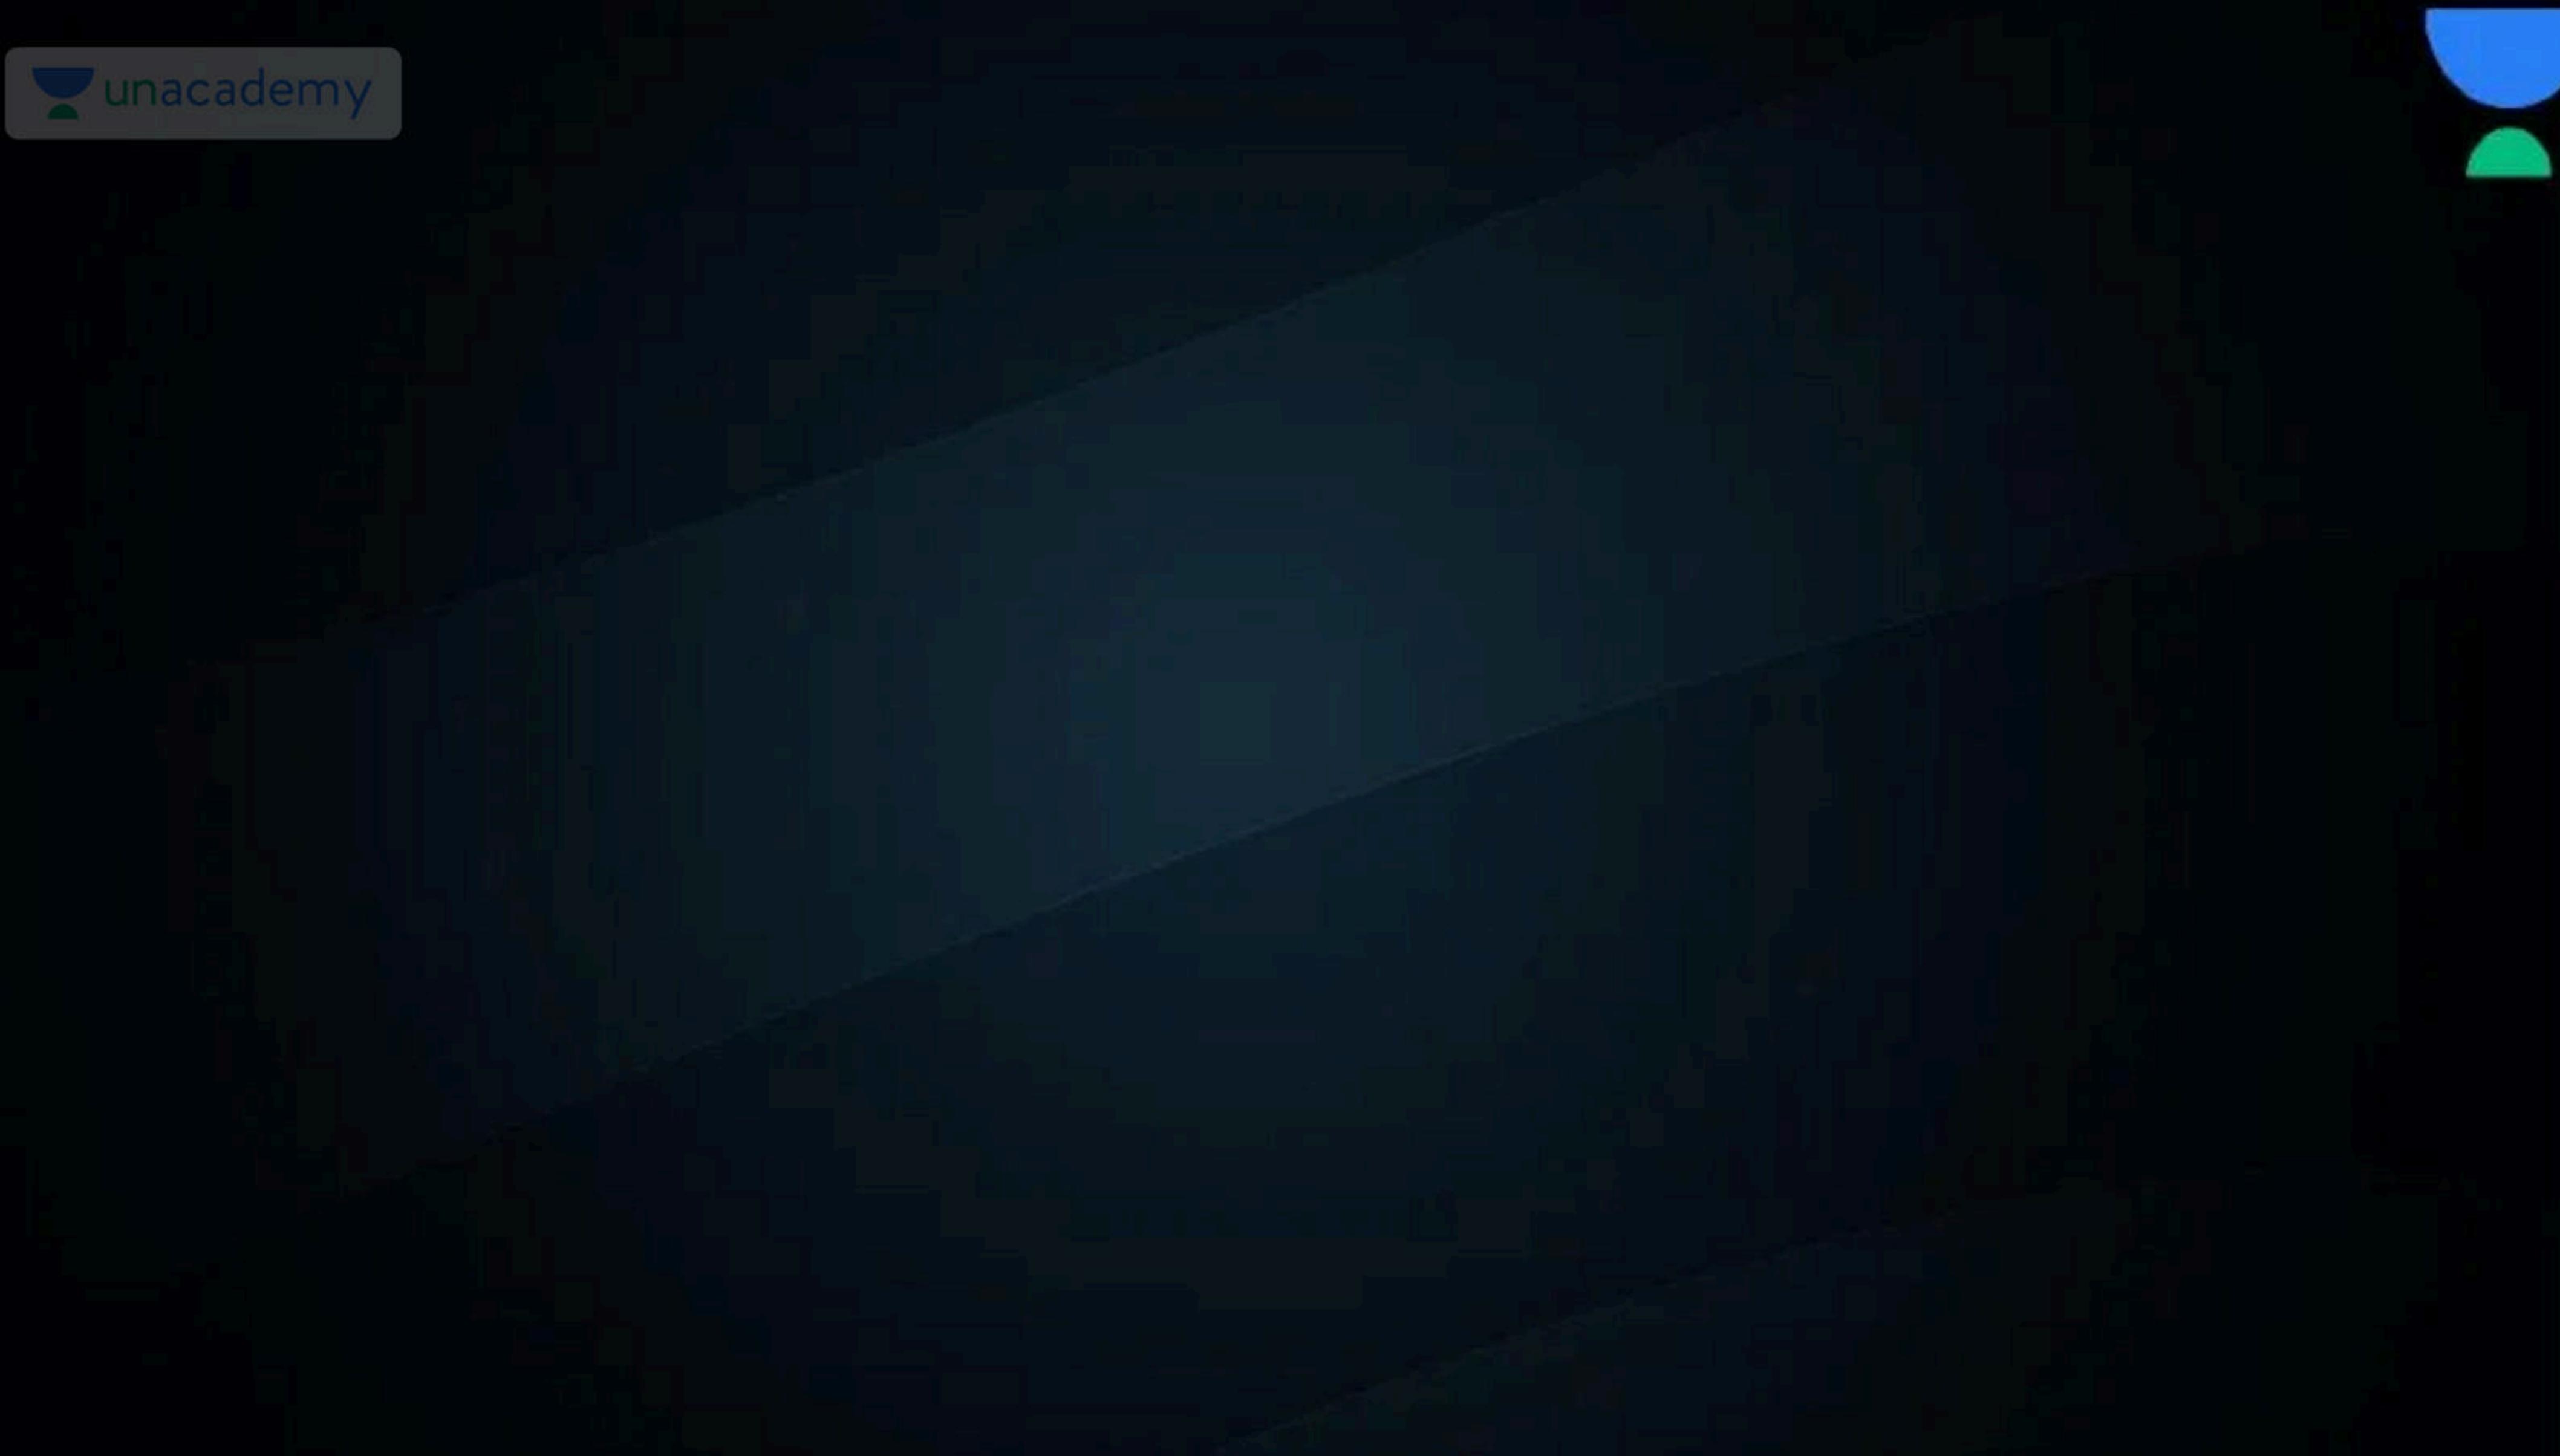
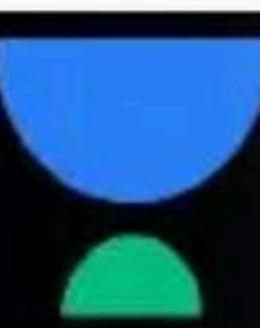
A Schedule is serializable Schedule if it is equivalent to a Serial Schedule.

- (i)Conflict Serializability
- (ii)View Serializability



Serializability

- Basic Assumption: Each transaction preserves database consistency.
- Thus, serial execution of a set of transactions preserves database consistency.
- A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule. Different forms of schedule equivalence give rise to the notions of:
 1. Conflict serializability
 2. view serializability



Conflict Serializability

- If a schedule S can be transformed into a schedule S' by a series of swaps of non-conflicting instructions, we say that S and S' are conflict equivalent.
- We say that a schedule S is conflict serializable if it is conflict equivalent to a serial schedule.

Conflict Serializability (Cont.)

- Schedule 3 can be transformed into Schedule 6, a serial schedule where T_2 follows T_1 , by series of swaps of non-conflicting instructions. Therefore Schedule 3 is conflict serializable.

Schedule 3

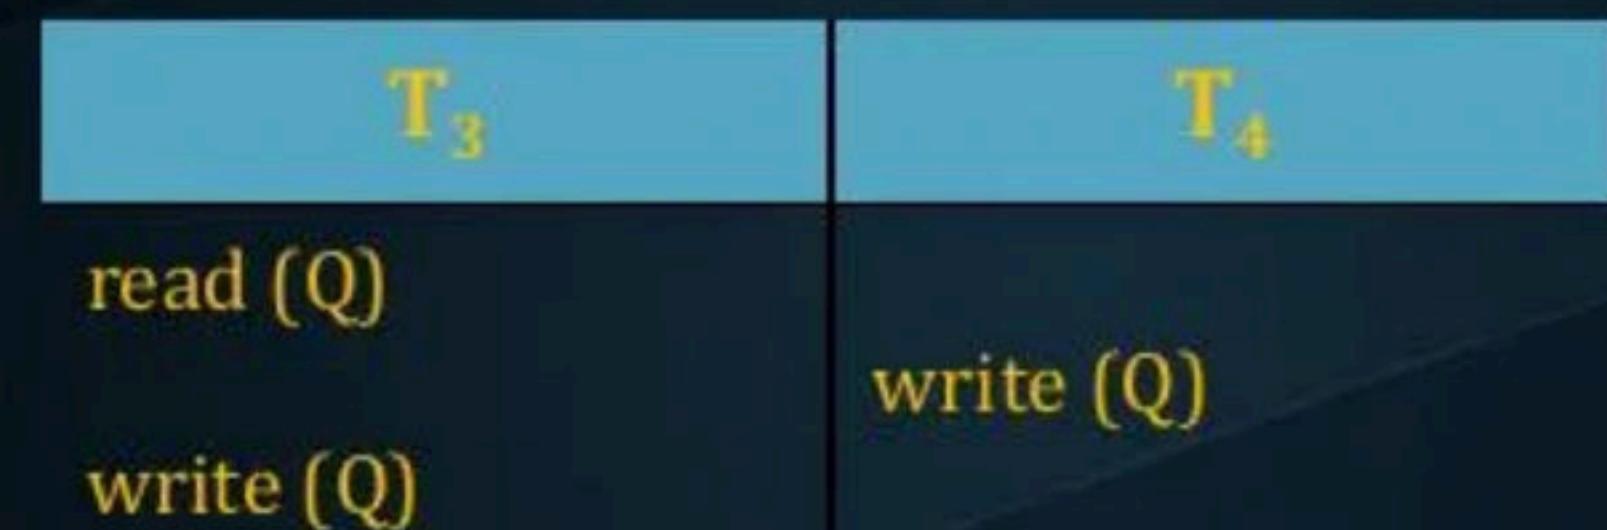
T_1	T_2
read (A)	
Write (A)	
	read (A)
	write (A)
read (B)	
write (B)	
	read (B)
	write (B)

Schedule 6

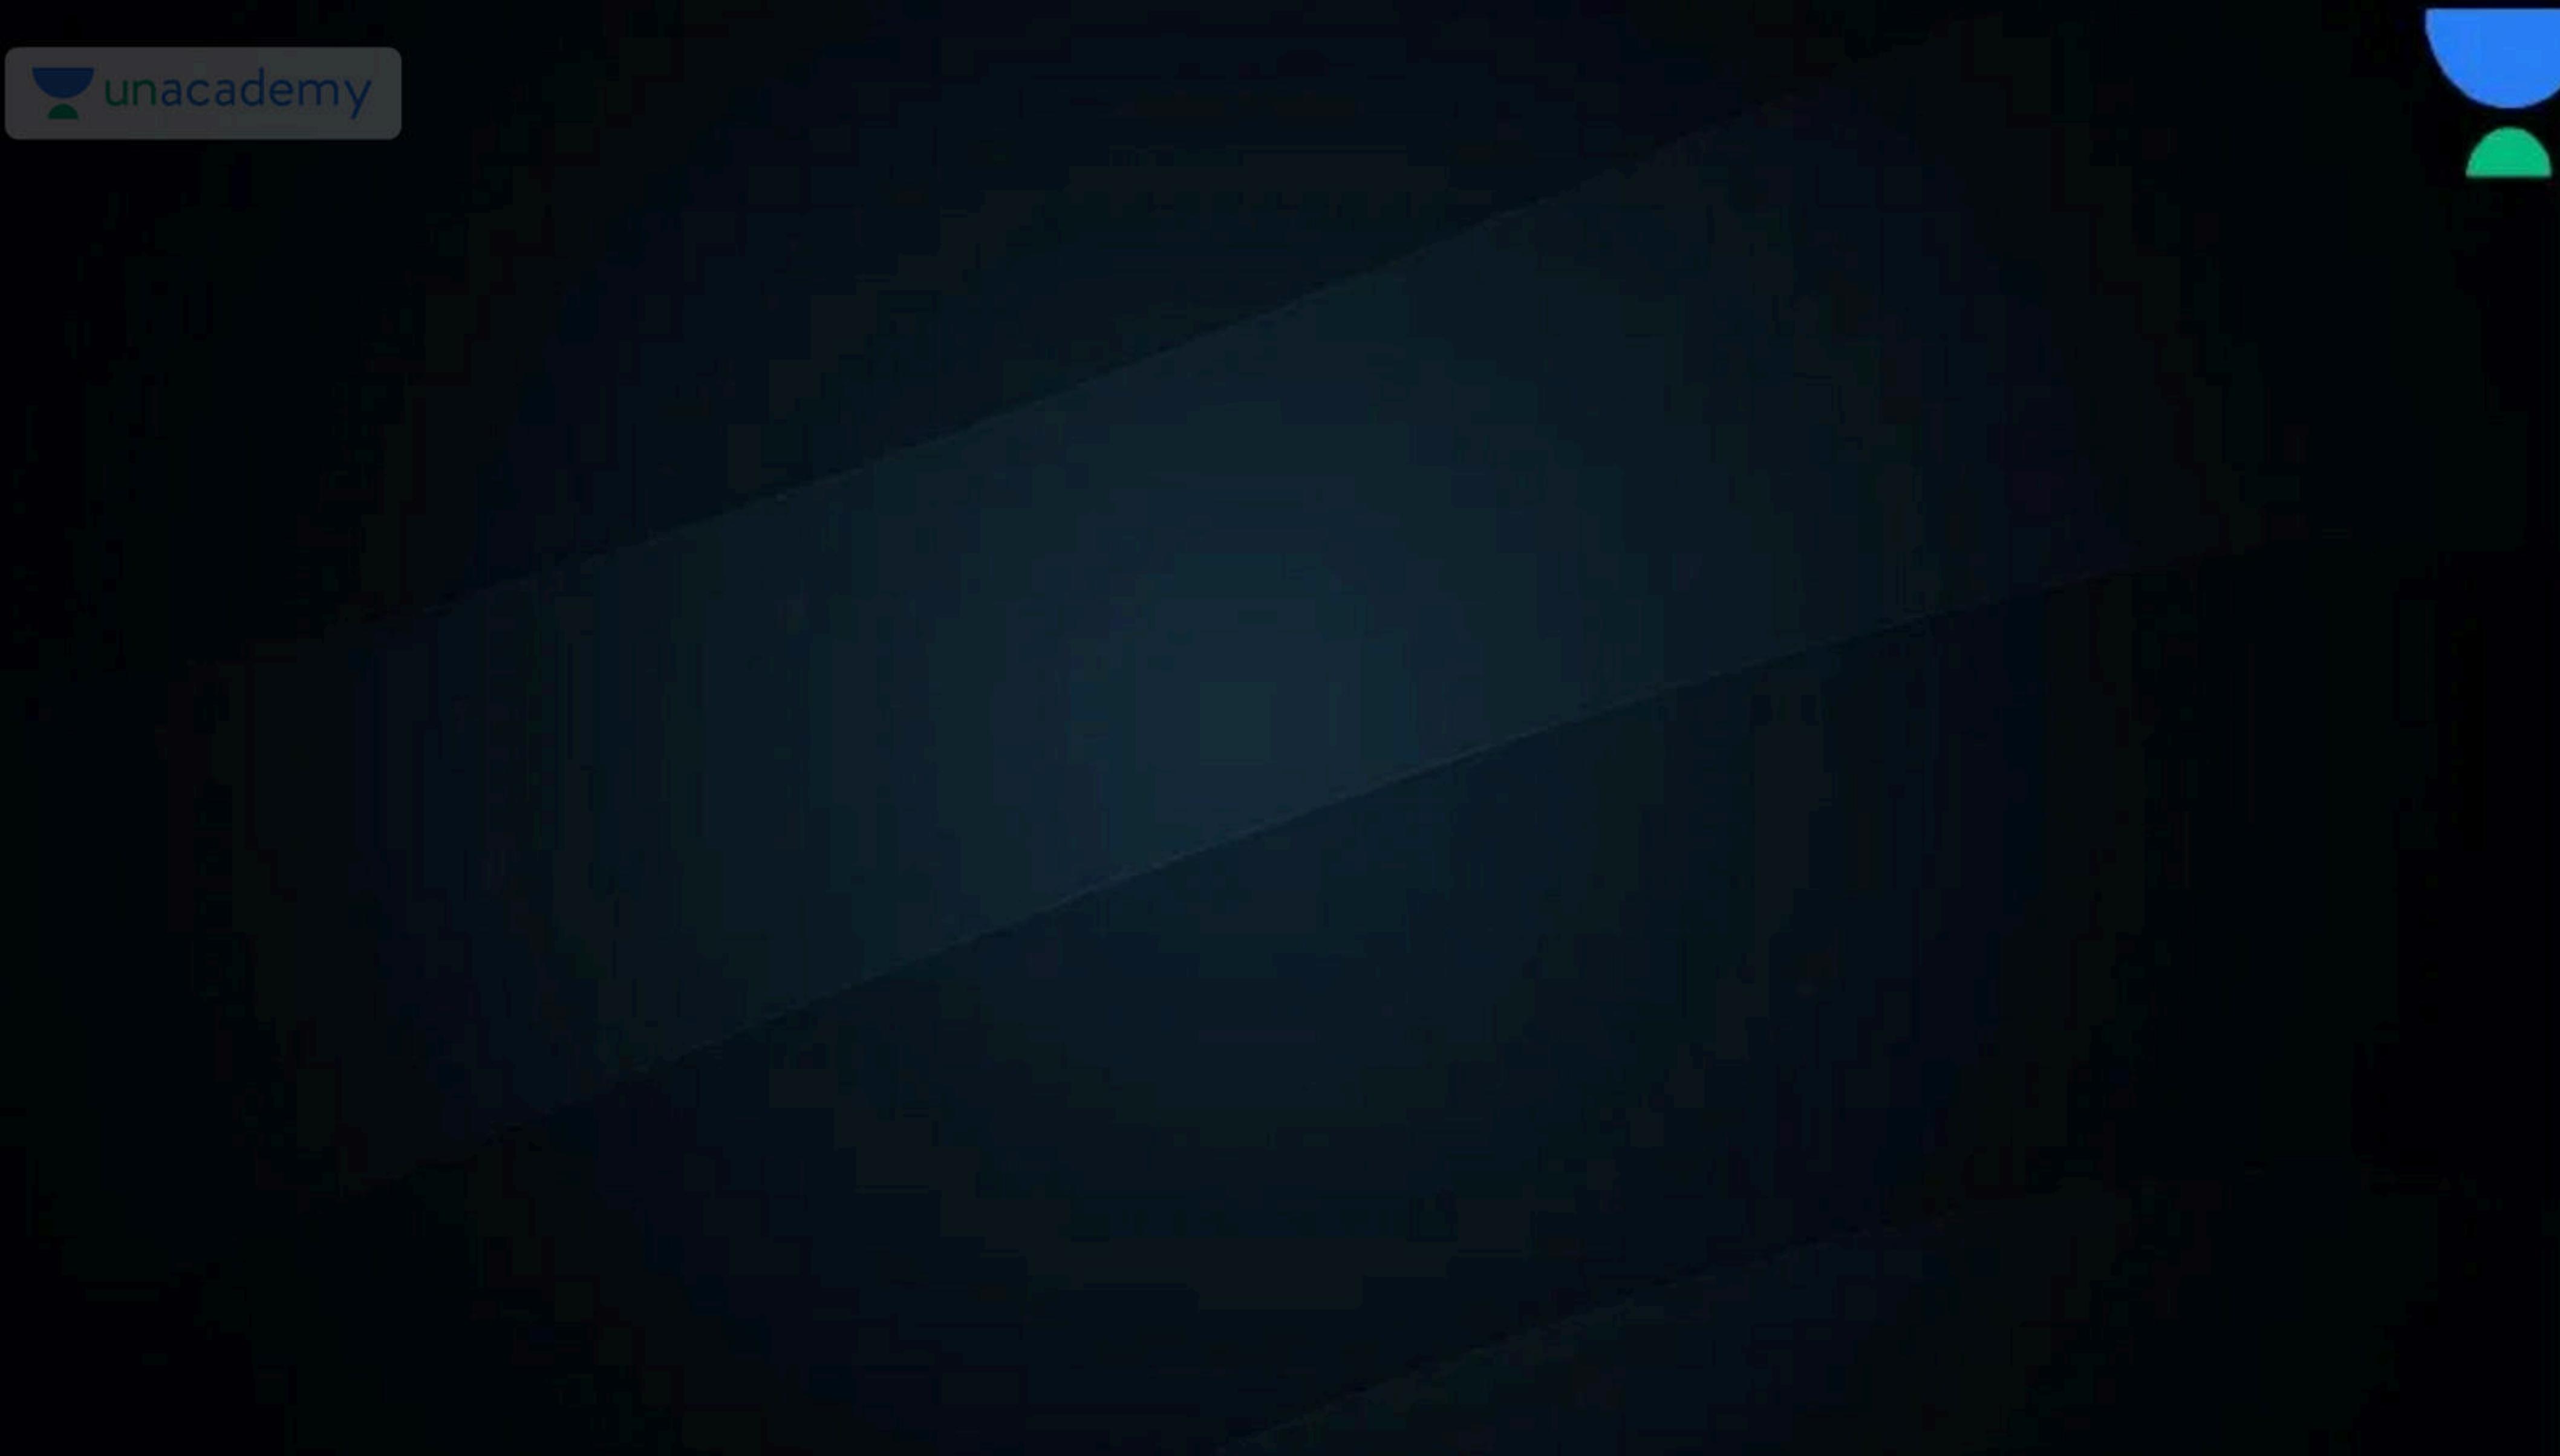
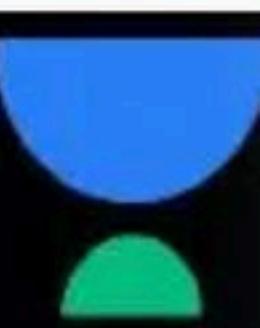
T_1	T_2
read (A)	
write (A)	
	read (B)
	write (B)
read (A)	
write (A)	
	read (B)
	write (B)

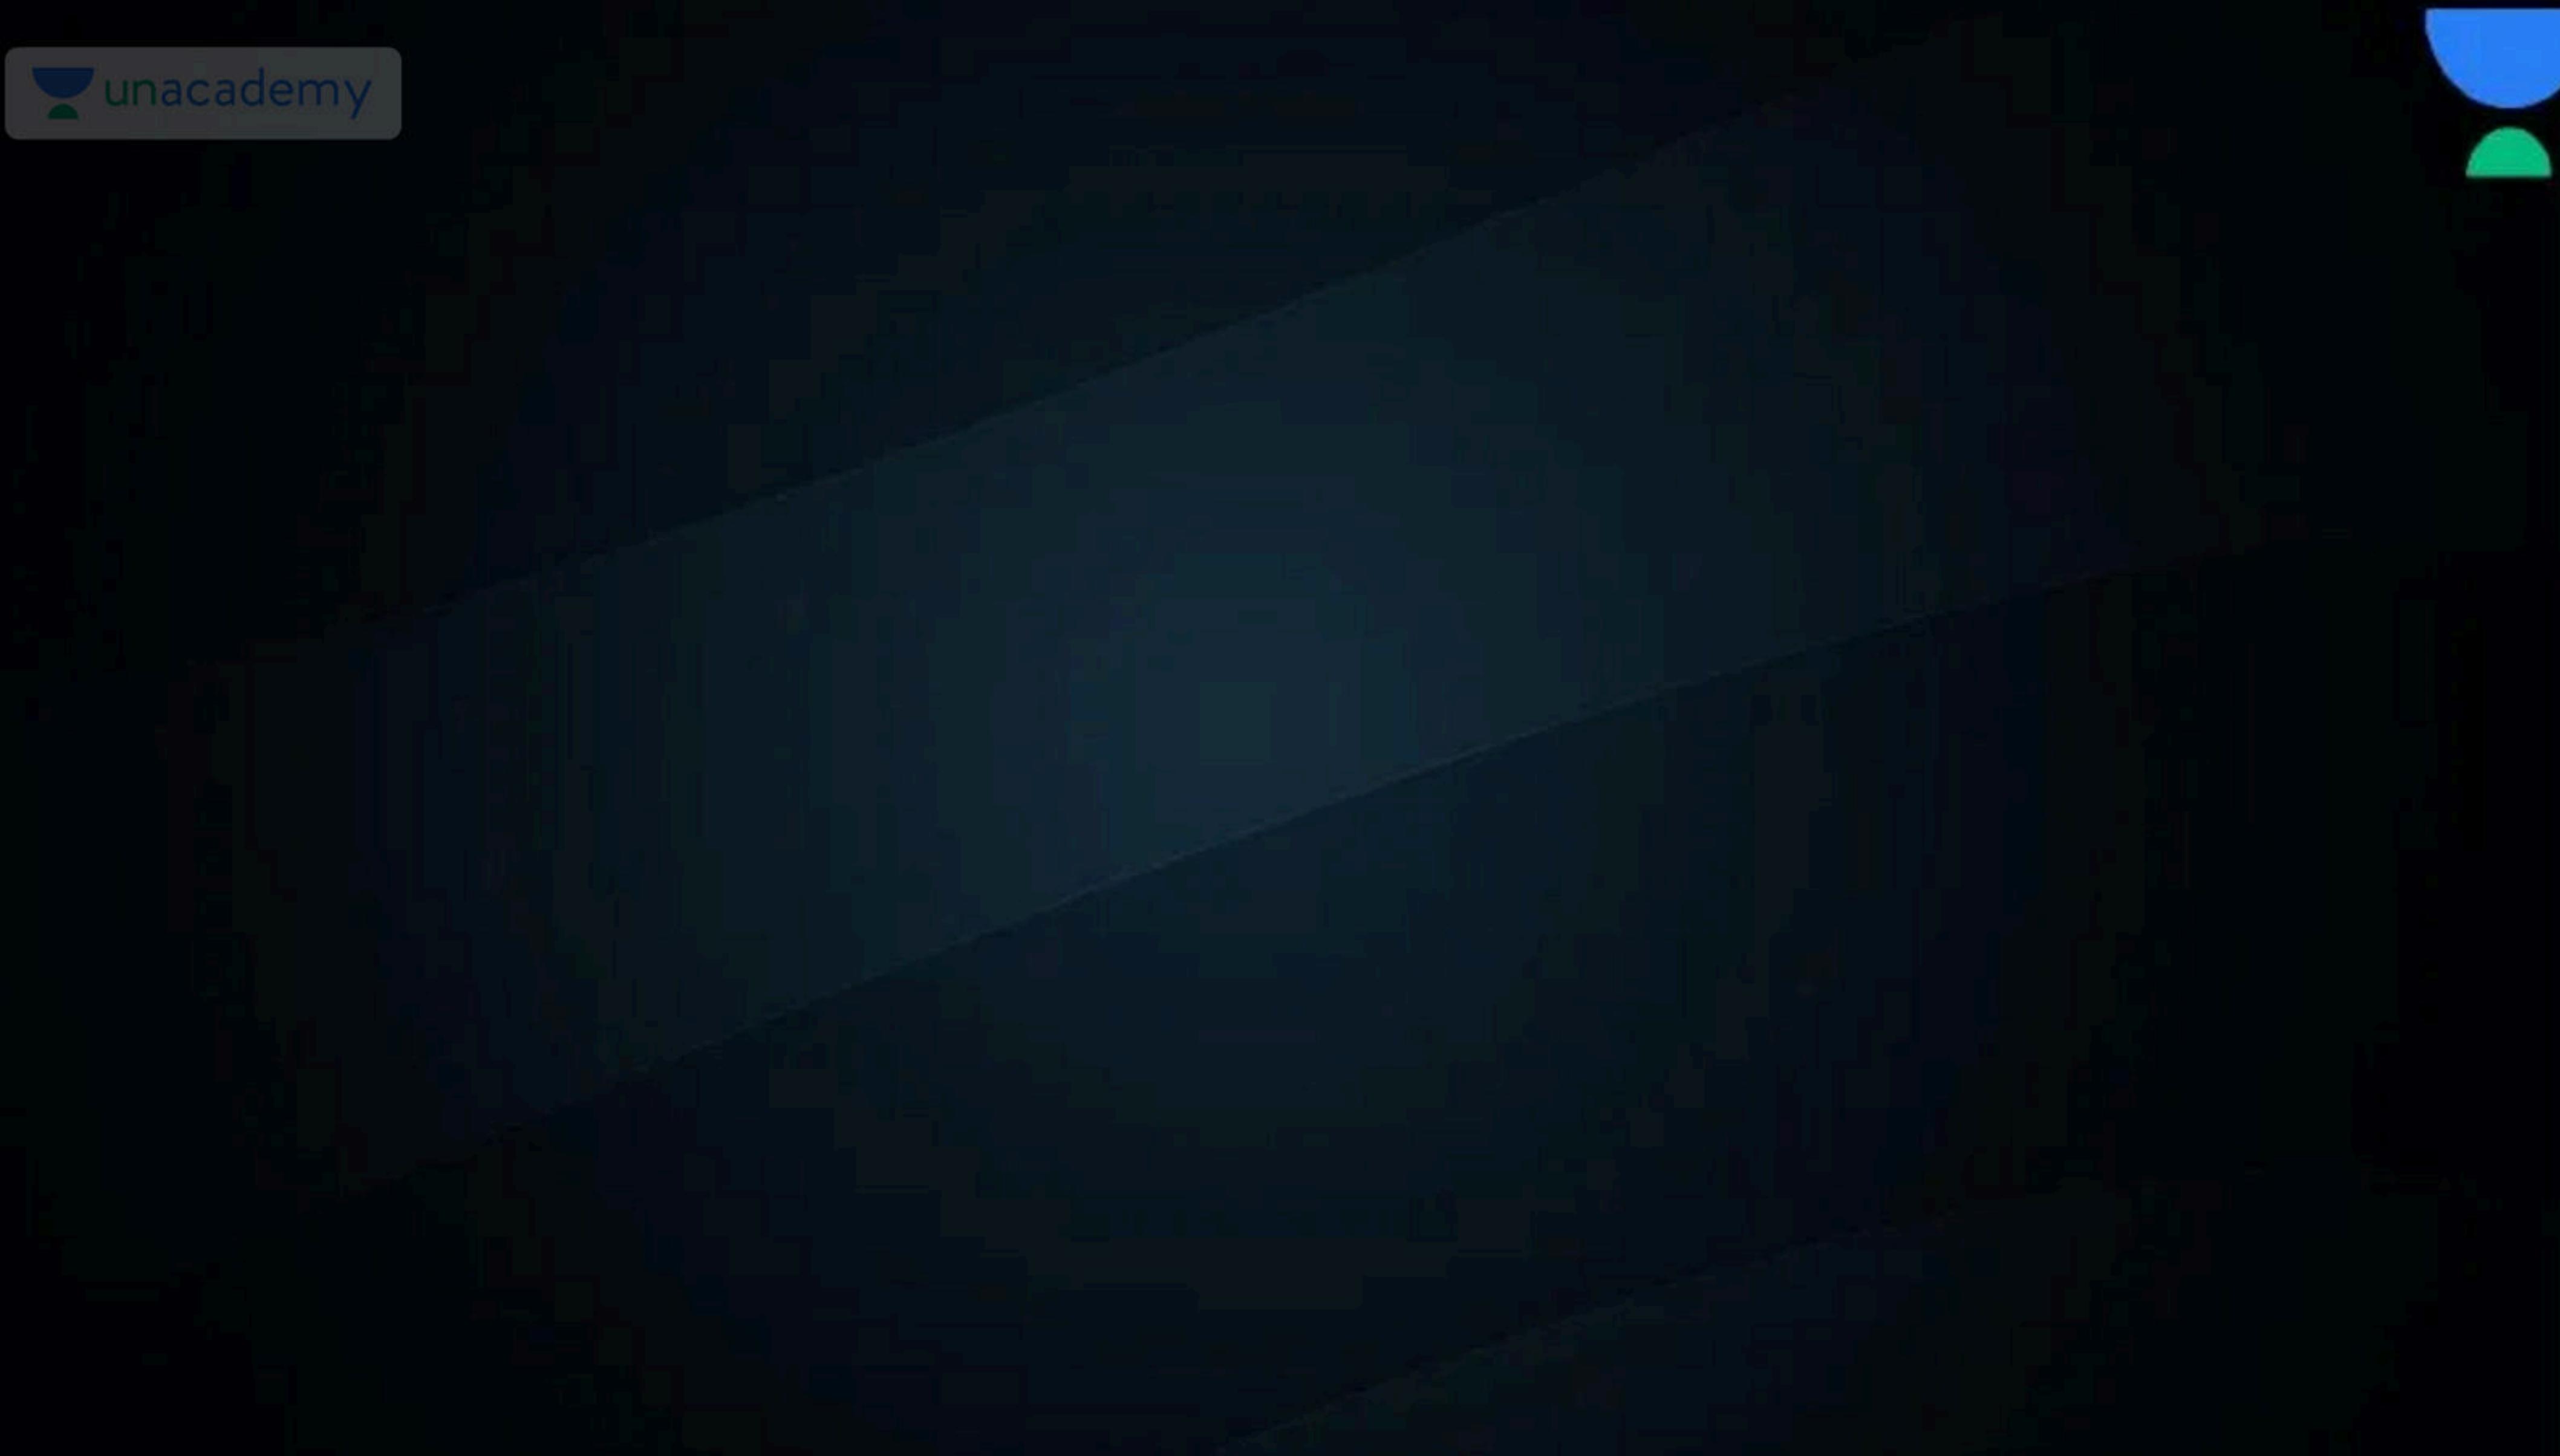
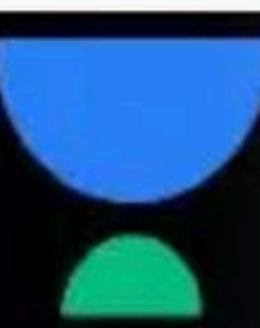
Conflict Serializability (Cont.)

- Example of a schedule that is not conflict serializable:



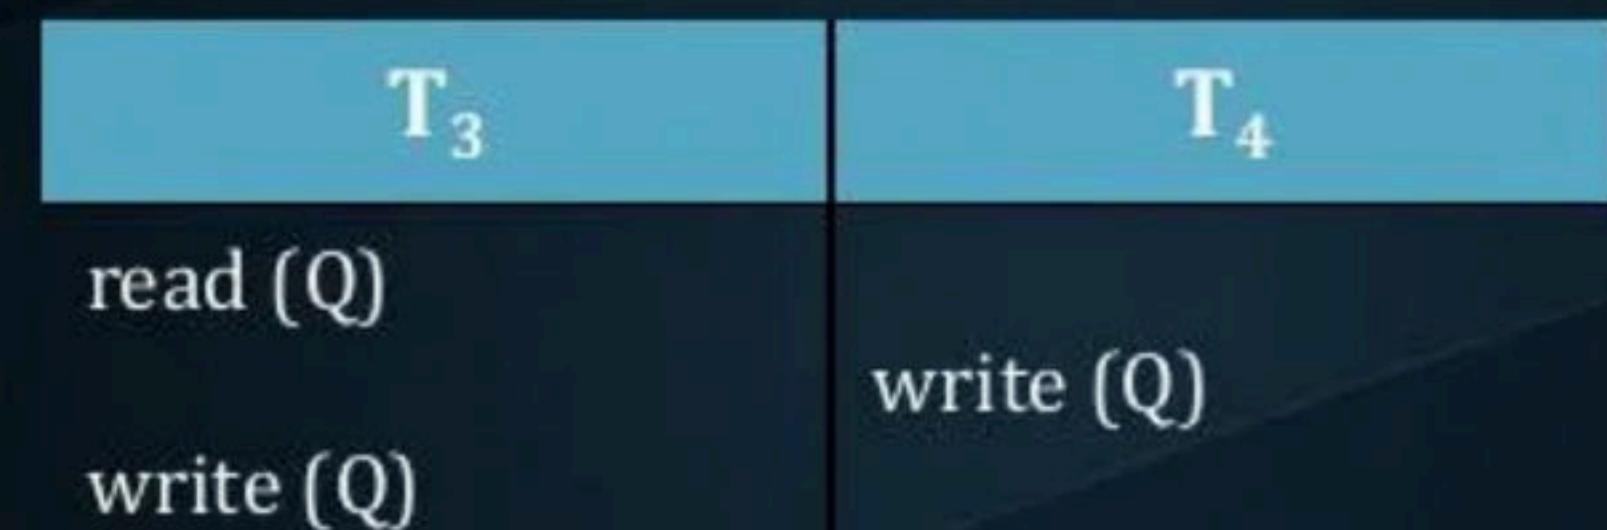
- We are unable to swap instructions in the above schedule to obtain either the serial schedule $\langle T_3, T_4 \rangle$, or the serial schedule $\langle T_4, T_3 \rangle$





Conflict Serializability (Cont.)

- Example of a schedule that is not conflict serializable:



- We are unable to swap instructions in the above schedule to obtain either the serial schedule $\langle T_3, T_4 \rangle$, or the serial schedule $\langle T_4, T_3 \rangle$

Conflict Serializable

A schedule is said to be conflict serializable if it is conflict equivalent to a serial schedule.

Same conflicting operation order in C_1 & S_1

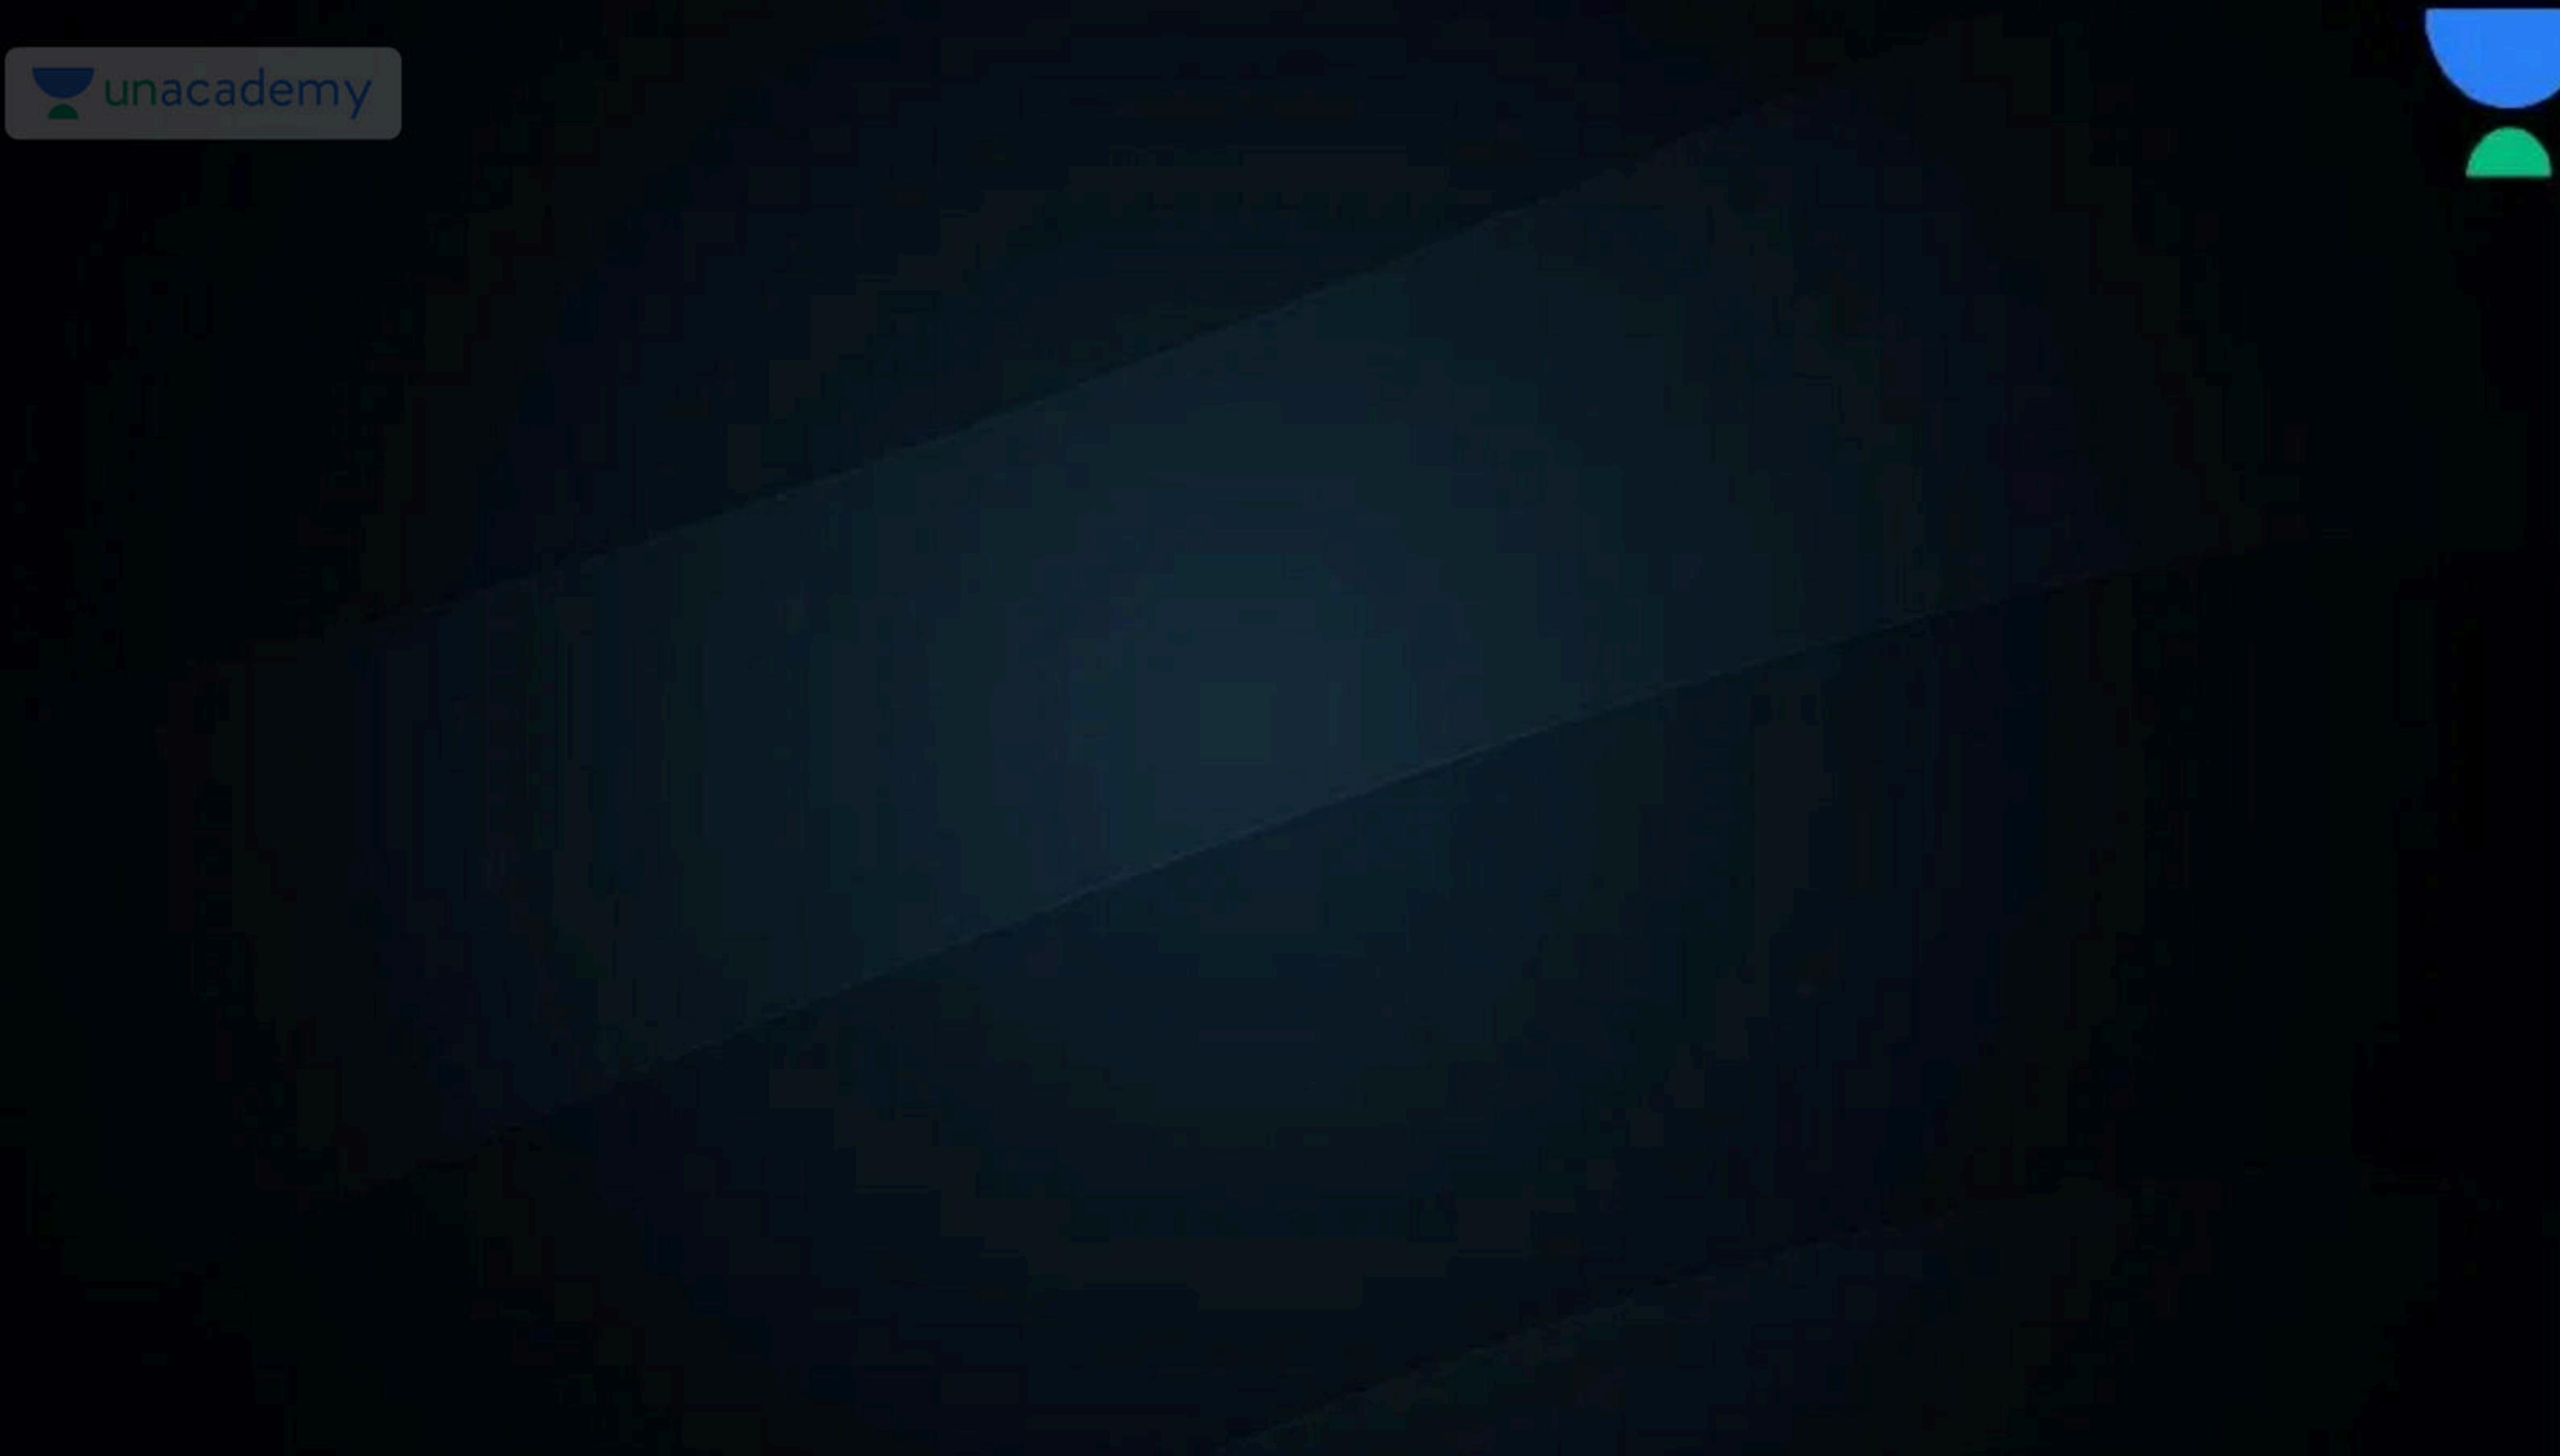
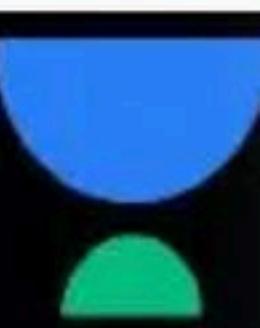
\therefore Its $\{C_1\}$ conflict is conflict serializable.

T_1	T_2
read(A)	
write(A)	
	read(A)
	write(A)
read(B)	
write(B)	
	read(B)
	write(B)

C_L

T_1	T_2
read(A)	
write(A)	
read(B)	
write(B)	
	read(A)
	write(A)
	read(B)
	write(B)

S_L

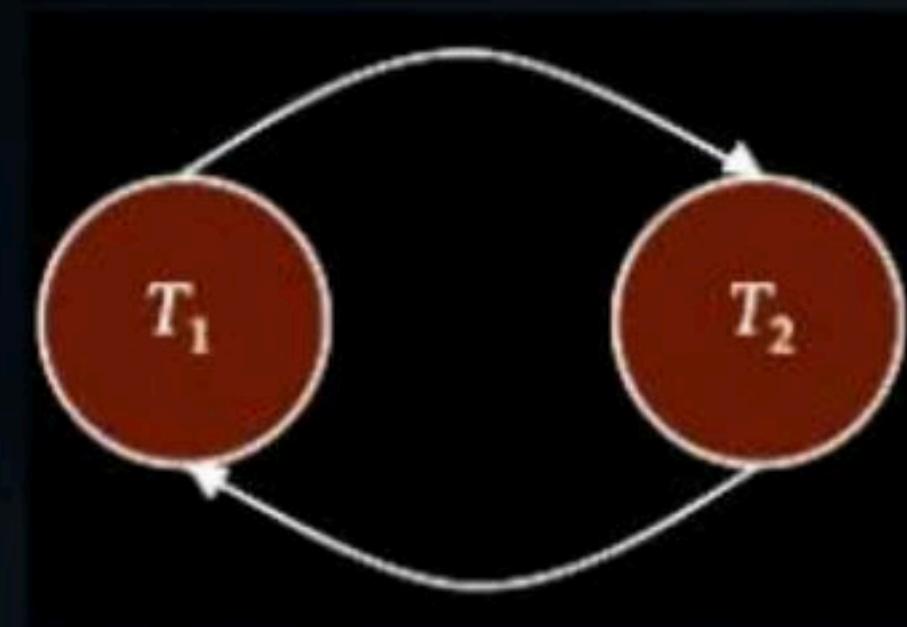


Conflicting Instructions

- Instructions l_i and l_j of transactions T_i and T_j respectively, conflict if and only if there exists some item Q accessed by both l_i and l_j , and at least one of these instructions wrote Q .
 1. $l_i = \text{read}(Q), l_j = \text{read}(Q)$. l_i and l_j don't conflict.
 2. $l_i = \text{read}(Q) l_j = \text{write}(Q)$. They conflict.
 3. $l_i = \text{write}(Q) l_j = \text{read}(Q)$. They conflict
 4. $l_i = \text{write}(Q) l_j = \text{write}(Q)$. They conflict
- Intuitively, a conflict between l_i and l_j forces a (logical) temporal order between them.
 - ❖ If l_i and l_j are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule.

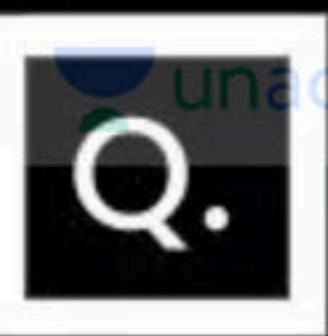
Testing for Serializability

- Testing for conflict serializability.
 - ❖ Consider some schedule of a set of transactions T_1, T_2, \dots, T_n
 - ❖ **Precedence graph** — a direct graph where the vertices are the transactions (names).
 - ❖ We draw an arc from T_i to T_j if the two transaction conflict, and T_i accessed the data item on which the conflict arose earlier.
 - ❖ We may label the arc by the item that was accessed.



A schedule is conflict serializable if and only if its precedence graph is acyclic.

NOTE: CNC [Cycle not conflict serializable]



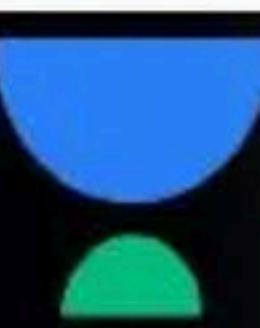
Q.

S: $R_1(A) W_1(A) R_2(A) W_2(A) R_1(B) W_1(B) R_2(B) W_2(B)$

T_1	T_2
$R(A)$ $W(A)$	$R(A)$ $W(A)$
$R(B)$ $W(B)$	$R(B)$ $W(B)$



R₁(A) R₂(A) W₂(A) W₁(A) R₁(B) W₁(B) R₂(B) W₂(B)



T ₁	T ₂
R(A)	
	R(A) W(A)
W(A) R(B) W(B)	
	R(B) W(B)



unacademy

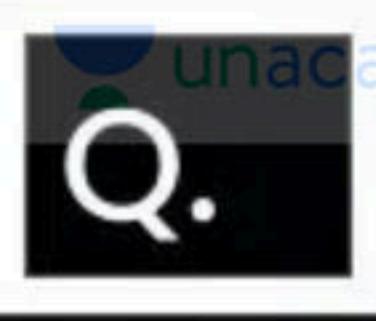
R₁(A) W₁(A) R₂(B) W₂(B) R₁(B) W₁(B) R₂(A) W₂(A)



Serializability Order

Important Point 1:

1. If S_1, S_2 Schedule are conflict equal then precedence graph of S_1 and S_2 must be same.
2. If S_1 and S_2 have same precedence graph then S_1 and S_2 may or may not conflict equal.



Q.

Consider the following schedules involving two transactions.
Which one of the following statements is TRUE?

S_1 : $r_1(X); r_1(Y); r_2(X); r_2(Y); w_2(Y); w_1(X)$

S_2 : $r_1(X); r_2(X); r_2(Y); W_2(Y); r_1(Y); w_1(X)$

[2007: 2 Marks]

- A Both S_1 and S_2 are conflict serializable
- B S_1 is conflict serializable and S_2 is not conflict serializable
- C S_1 is not conflict serializable and S_2 is conflict serializable
- D Both S_1 and S_2 are not conflict serializable

Consider the following four schedules due to three transactions (indicated by the subscript) using read and write on a data item x , denoted by $r(x)$ and $w(x)$ respectively. Which one of them is conflict serializable?

[2014(Set-1): 2 Marks]

- A $r_1(x); r_2(x); w_1(x); r_3(x); w_2(x)$
- B $r_2(x); r_1(x); w_2(x); r_3(x); w_1(x)$
- C $r_3(x); r_2(x); r_1(x); w_2(x); w_1(x)$
- D $r_2(x); w_2(x); r_3(x); r_1(x); w_1(x)$



Let $r_i(z)$ and $w_i(z)$ denote read and write operations respectively on a data item by a transaction T_i . Consider the following two schedules.

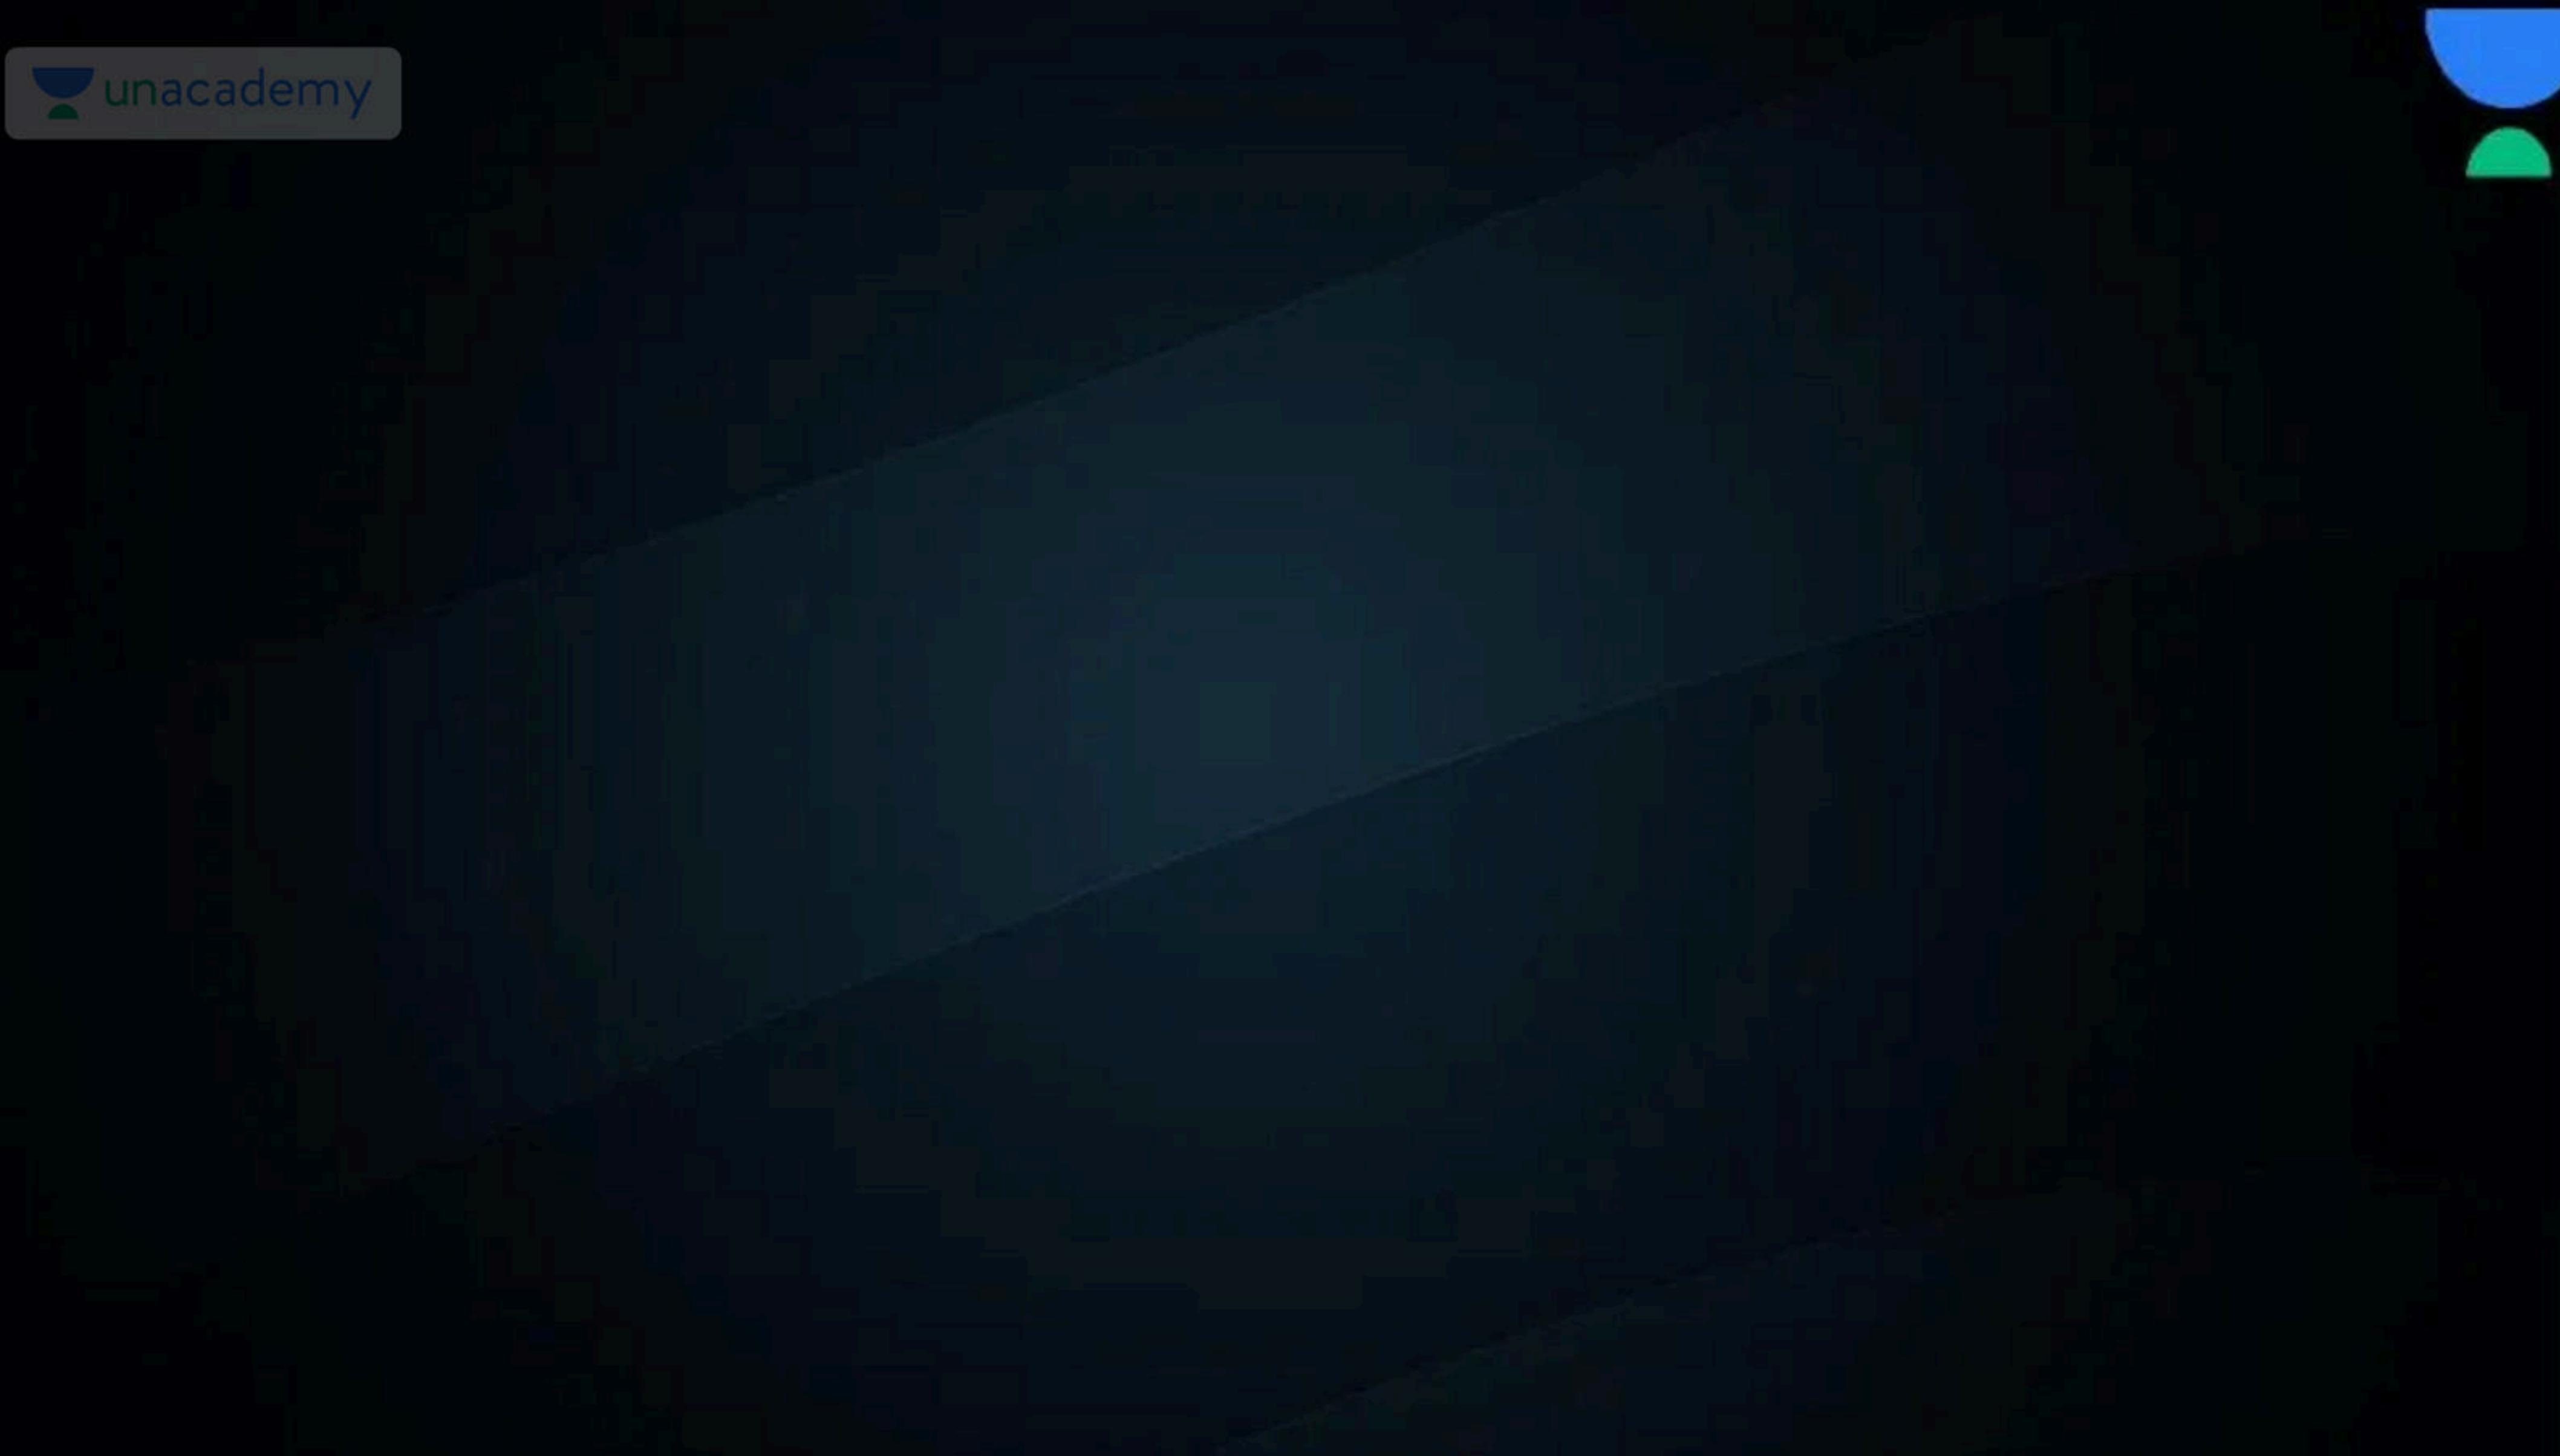
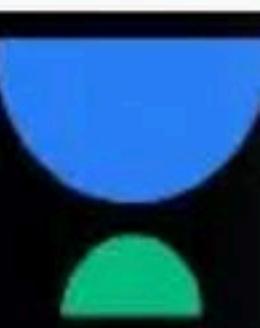
$S_1: r_1(x) \ r_1(y) \ r_2(x) \ r_2(y) \ w_2(y) \ w_1(x)$

$S_2: r_1(x) \ r_2(x) \ r_2(y) \ w_2(y) \ r_1(y) \ w_1(x)$

Which one of the following options is correct?

[MCQ: 2021: 2M]

- A** S_1 is conflict serializable, and S_2 is not conflict serializable.
- B** S_1 is not conflict serializable, and S_2 is conflict serializable.
- C** Both S_1 and S_2 are conflict serializable.
- D** Neither S_1 nor S_2 is conflict serializable.



Q.1

unacademy

Consider the following schedules involving two transactions.
Which one of the following statements is TRUE?

$S_1: r_1(X); r_1(Y); r_2(X); r_2(Y); w_2(Y); w_1(X)$

$S_2: r_1(X); r_2(X); r_2(Y); W_2(Y); r_1(Y); w_1(X)$

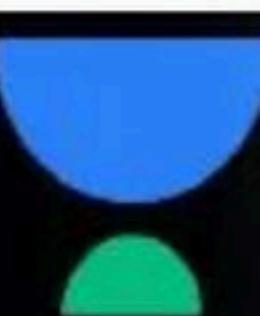
[2007: 2 Marks]

- A Both S_1 and S_2 are conflict serializable
- B S_1 is conflict serializable and S_2 is not conflict serializable
- C S_1 is not conflict serializable and S_2 is conflict serializable
- D Both S_1 and S_2 are not conflict serializable

Consider the following four schedules due to three transactions (indicated by the subscript) using read and write on a data item x , denoted by $r(x)$ and $w(x)$ respectively. Which one of them is conflict serializable?

[2014(Set-1): 2 Marks]

- A $r_1(x); r_2(x); w_1(x); r_3(x); w_2(x)$
- B $r_2(x); r_1(x); w_2(x); r_3(x); w_1(x)$
- C $r_3(x); r_2(x); r_1(x); w_2(x); w_1(x)$
- D $r_2(x); w_2(x); r_3(x); r_1(x); w_1(x)$



Consider the transactions T1, T2 and T3 and the schedules S1 and S2 given below.

T1: r1(X); r1(Z); w1(X); w1(Z)

T2: r2(Y) ; r2(Z) ; w2(Z)

T3: r3(Y); r3(X); w3(Y)

S1: r1(X); r3(Y); r3(X); r2(Y); r2(Z); w3(Y); w2(Z); r1(Z); w1(X); w1(Z)

S2: r1(X); r3(Y); r2(Y); r3(X); r1(Z); r2(Z); w3(Y); w1(X); w2(Z); w1(Z)

Which one of the following statements about the schedules is TRUE?

[GATE-2014-CS: 2M]

- A Only S1 is conflict-serializable.
- B Only S2 is conflict-serializable.
- C Both S1 and S2 are conflict-serializable.
- D Neither S1 nor S2 is conflict-serializable.

Q.4

Let $r_i(z)$ and $w_i(z)$ denote read and write operations respectively on a data item by a transaction T_i . Consider the following two schedules.

$S_1: r_1(x) \ r_1(y) \ r_2(x) \ r_2(y) \ w_2(y) \ w_1(x)$

$S_2: r_1(x) \ r_2(x) \ r_2(y) \ w_2(y) \ r_1(y) \ w_1(x)$

Which one of the following options is correct?

[MCQ: 2021: 2M]

- A S_1 is conflict serializable, and S_2 is not conflict serializable.
- B S_1 is not conflict serializable, and S_2 is conflict serializable.
- C Both S_1 and S_2 are conflict serializable.
- D Neither S_1 nor S_2 is conflict serializable.

Let $R_i(z)$ and $W_i(z)$ denote read and write operations on a data element z by a transaction T_i , respectively. Consider the schedule S with four transactions.

$S: R_4(x), R_2(x), R_3(x), R_1(y), W_1(y), W_2(x), W_3(y), R_4(y)$

Which one of the following serial schedules is conflict equivalent to S ?

[2022: 2 Marks]

A $T_1 \rightarrow T_3 \rightarrow T_4 \rightarrow T_2$

B $T_1 \rightarrow T_4 \rightarrow T_3 \rightarrow T_2$

C $T_4 \rightarrow T_1 \rightarrow T_3 \rightarrow T_2$

D $T_3 \rightarrow T_1 \rightarrow T_4 \rightarrow T_2$

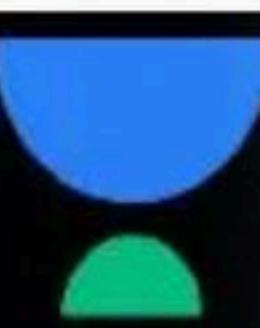
Consider the following transaction involving two bank accounts x and y.

read(x); x: = x - 50; write (x); read (y); y: = y + 50; write (y)

The constraint that the sum of the accounts x and y should remain constant is that of

[2015(Set-2): 1 Marks]

- A Atomicity
- B Consistency
- C Isolation
- D Durability



Which one of the following is NOT a part of the ACID properties of database transactions?

[GATE-2016-CS: 1M]

- A Atomicity
- B Consistency
- C Isolation
- D Deadlock-freedom

Suppose a database schedule S involves transaction T_1, \dots, T_n . Construct the precedence graph of S with vertices representing the transactions and edges representing the conflicts. If S is serializable, which one of the following orderings of the vertices of the precedence graph is guaranteed to yield a serial schedule?

[GATE-2016-CS: 2M]

- A Topological order
- B Depth-first order
- C Breadth-first order
- D Ascending order of transaction indices

Consider the following schedule for transactions T1, T2 and T3:

Which one of the schedules below is the correct serialization of the above?

[GATE-2010-CS: 2M]

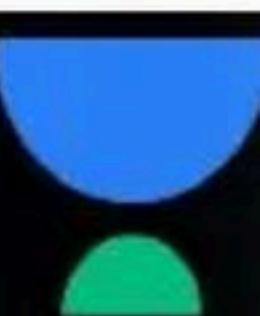
T1	T3	T3
Read(X)		
	Read(Y)	
		Read(Y)
	Write(Y)	
Write(X)		
		Write(X)
	Read(X)	
	Write(X)	

A $T_1 \rightarrow T_3 \rightarrow T_2$

C $T_2 \rightarrow T_3 \rightarrow T_1$

B $T_2 \rightarrow T_1 \rightarrow T_3$

D $T_3 \rightarrow T_1 \rightarrow T_2$



Consider two transactions T_1 and T_2 , and four schedules S_1, S_2, S_3, S_4 of T_1 and T_2 as given below:

$T_1: R_1[x] W_1[x] W_1[y];$

$T_2: R_2[x] R_2[y] W_2[y];$

$S_1: R_1[x] R_2[x] R_2[y] W_1[x] W_1[y] W_2[y];$

$S_2: R_1[x] R_2[x] R_2[y] W_1[x] W_2[y] W_1[y];$

$S_3: R_1[x] W_1[x] R_2[x] W_1[y] R_2[y] W_2[y];$

$S_4: R_2[x] R_2[y] R_1[x] W_1[x] W_1[y] W_2[y];$

Which of the above schedules are conflict serializable?

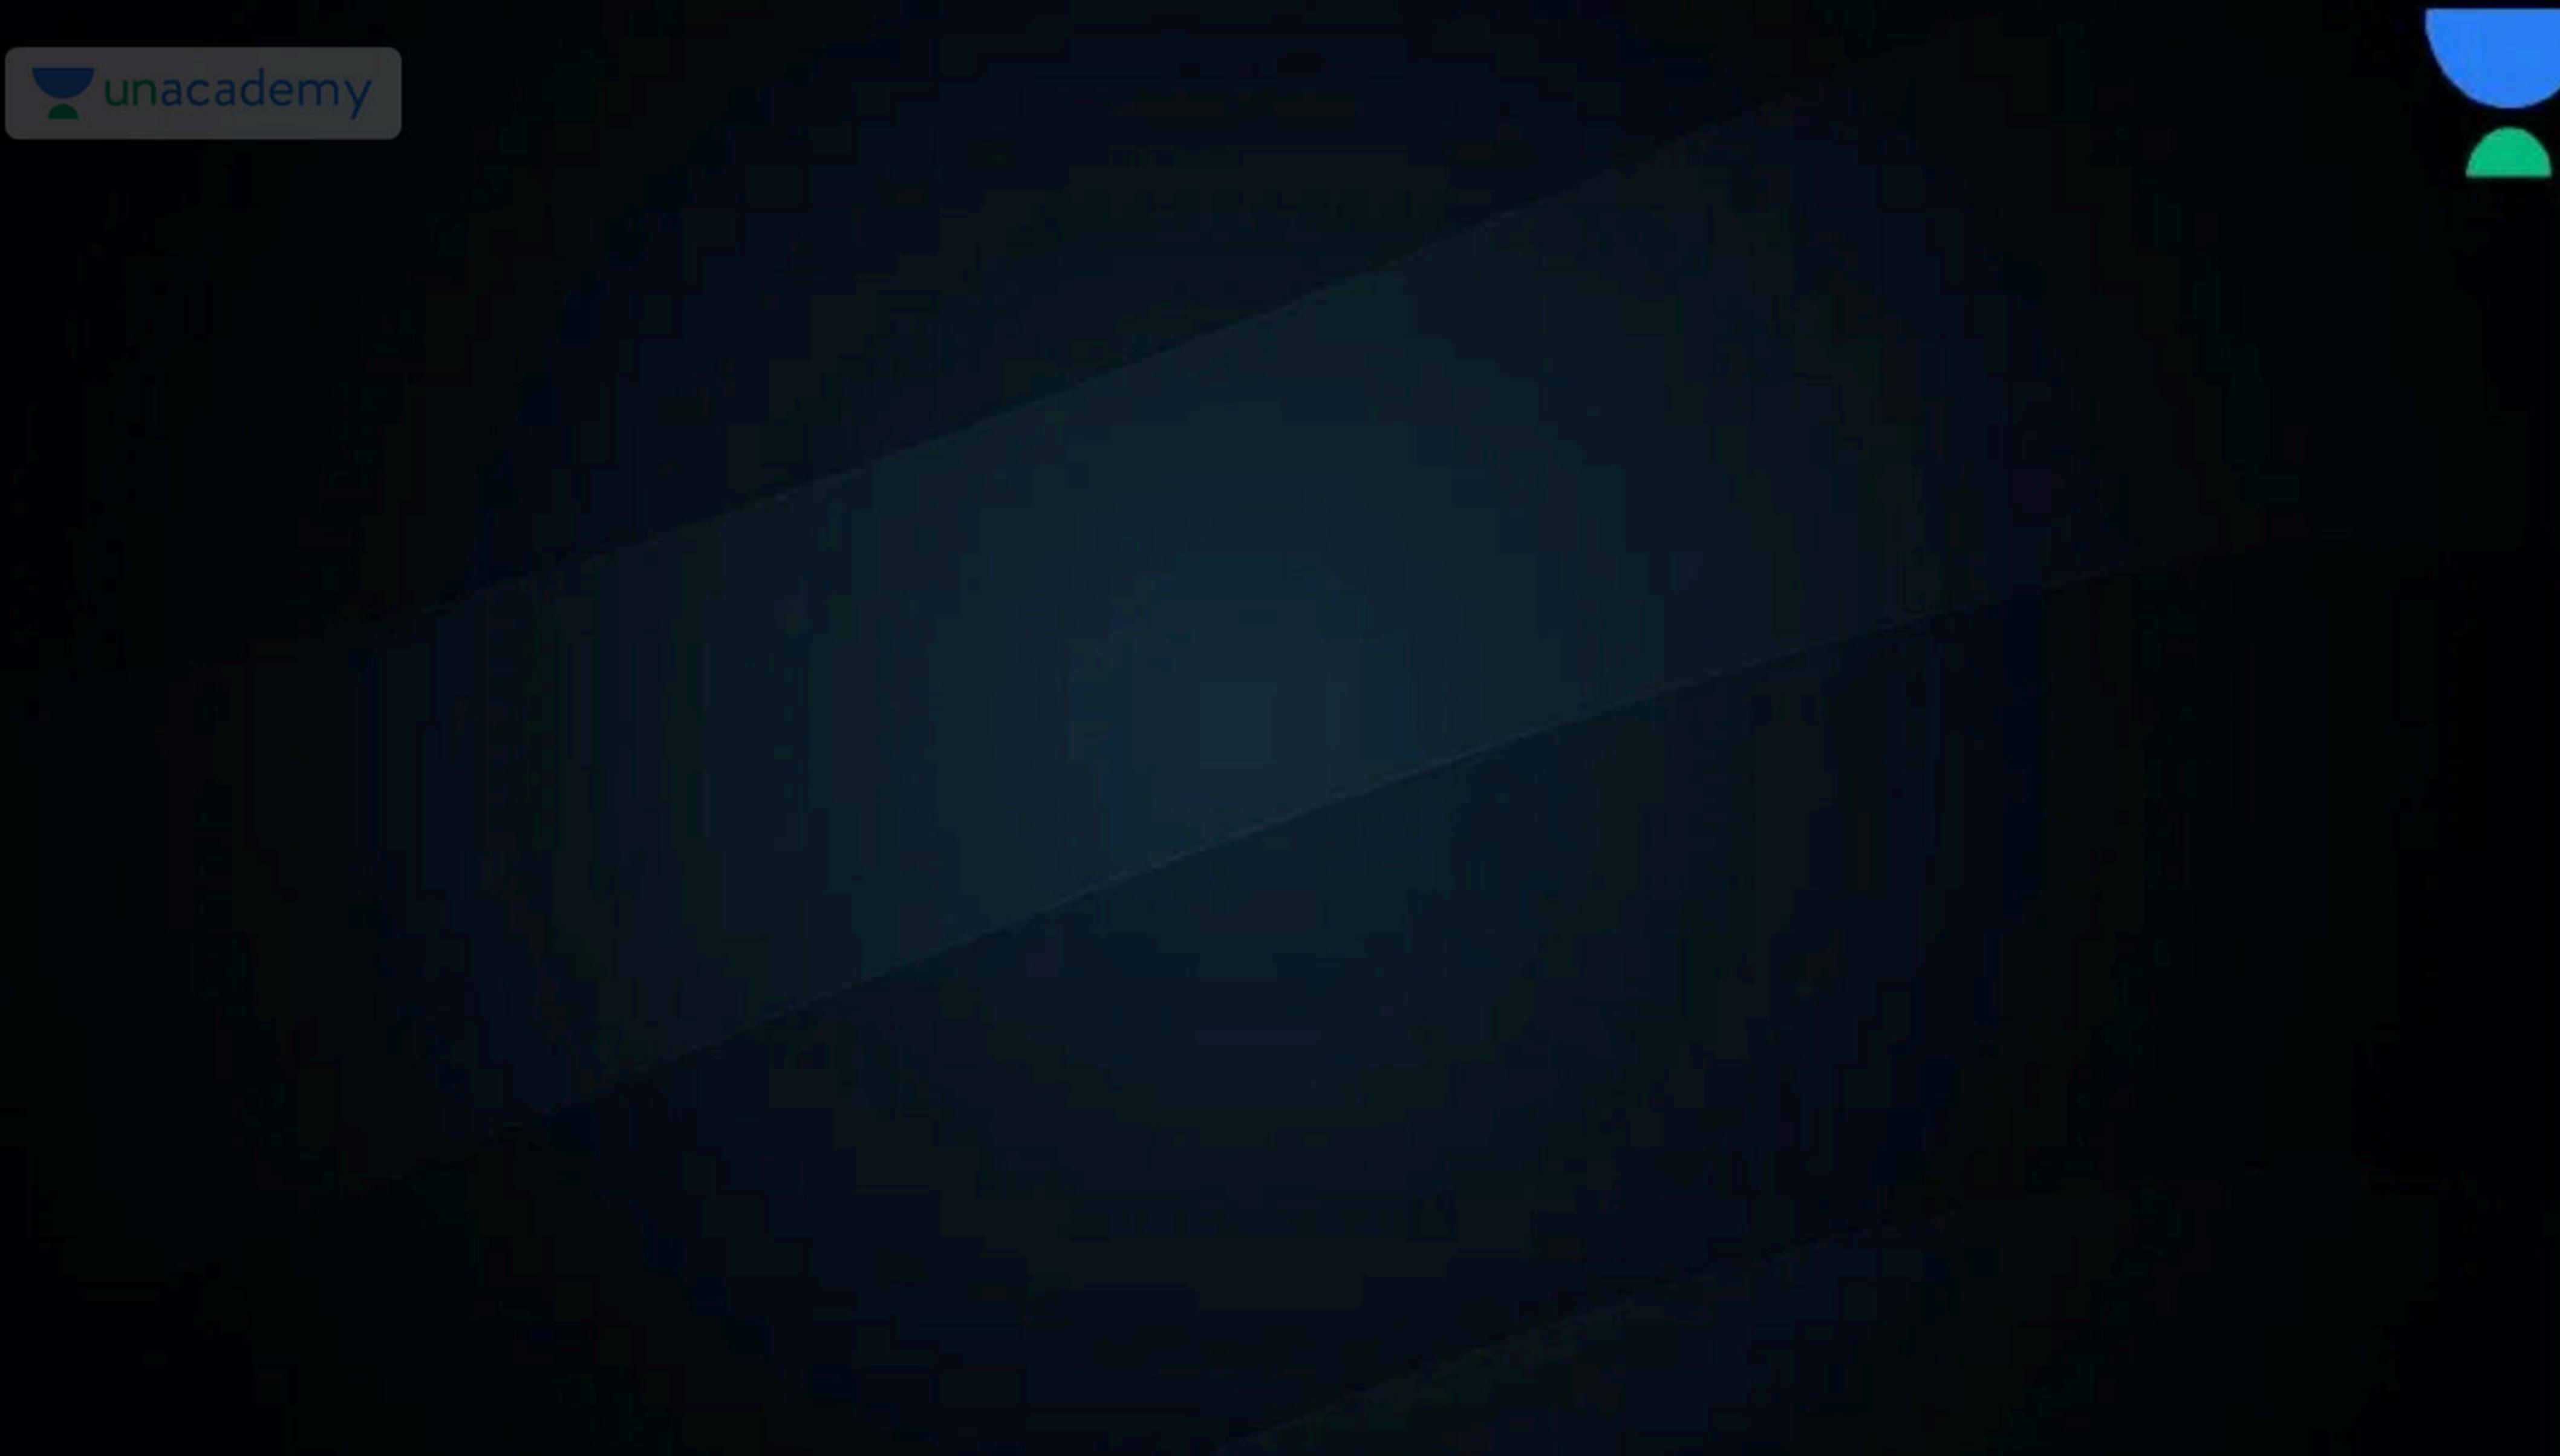
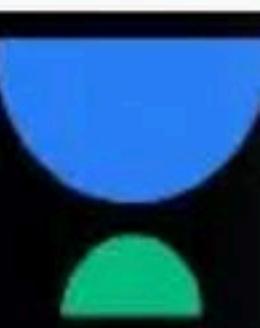
[GATE-2009-CS: 2M]

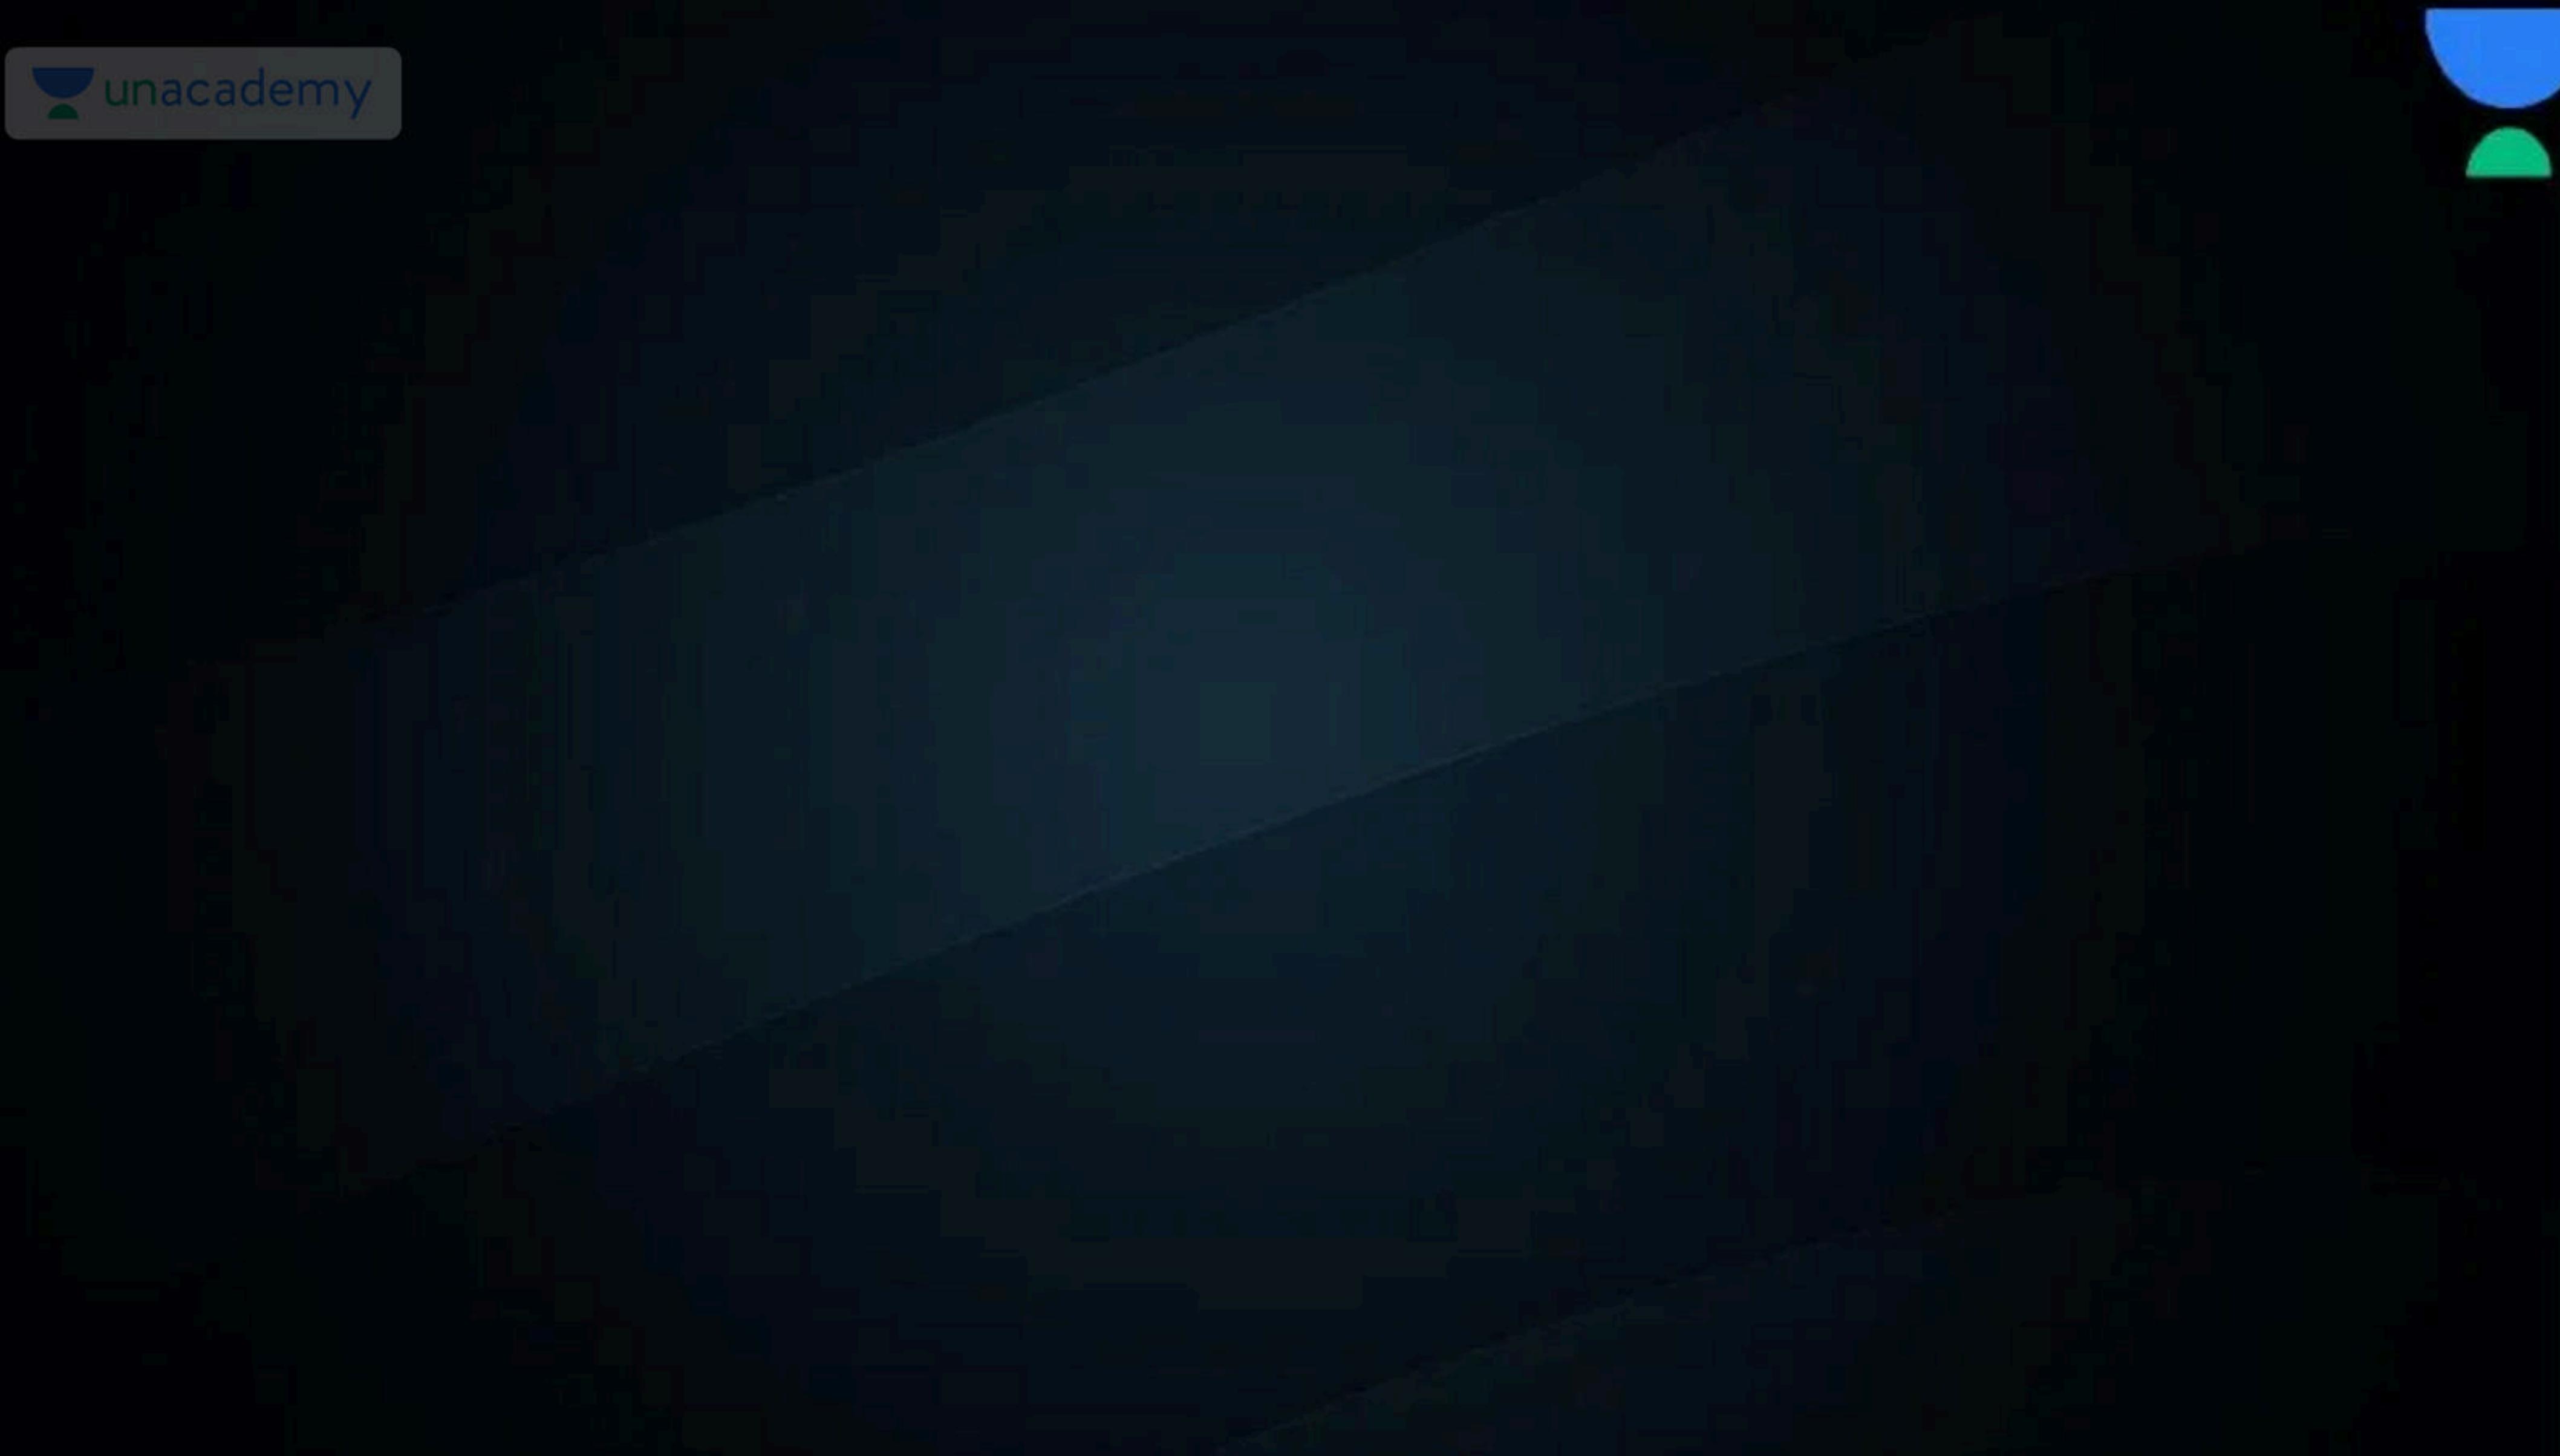
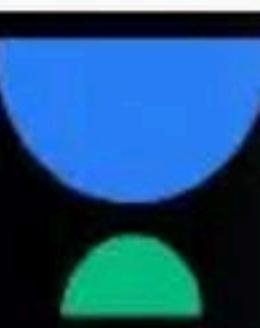
A S_1 and S_2

B S_2 and S_3

C S_3 only

D S_4 only







THANK YOU!

Here's to a cracking journey ahead!